



Faculty of Science



Introduction: Hardware Trends and List Homomorphism

Cosmin E. Oancea and Troels Henriksen
`[cosmin.oancea,athas]@diku.dk`

Department of Computer Science (DIKU)
University of Copenhagen

September 2015 PMPH Lecture Notes



1 Introduction²

- Brief History
- Computer Architecture Definition
- Course Organization

2 Trends of Critical Components of a Parallel System

- Processor
- Memory
- Interconnect

3 Technological Challenges/Constraints

- Power
- Reliability
- Wire Delays
- Design Complexity
- CMOS Endpoint

4 How to Reason about Parallelism?

- Amdahl's Law
- List Homomorphism
- List-Homomorphism \equiv Map-Reduce



Introduction²

- Past > 20 years Information Revolution:
Explosive growth of semiconductor integration + Internet.
- Moore's Law 1960s:
 - computing power doubles every 19-24 months
 - system cost effectiveness = performance/cost grows exp.
 - CMOS endpoint is near: miniaturization reaches its limits (complementary metal-oxide semiconductor).



Introduction²

- Past > 20 years **Information Revolution**:
Explosive growth of semiconductor integration + Internet.
- Moore's Law 1960s:
 - computing power doubles every 19-24 months
 - system cost effectiveness = performance/cost grows exp.
 - **CMOS endpoint is near**: miniaturization reaches its limits (complementary metal-oxide semiconductor).
- Improved Chip Design:
 - each new process generation \Rightarrow higher clock rates
 - logic switching speed & amount of on-chip logic increase (\uparrow)
 - better circuit design & deep pipelines \Rightarrow fewer gate delays / stage
 - \uparrow on-chip resources \Rightarrow \uparrow throughput by parallelism at all stages.
- Computer Architecture goes back to before 1970s:
How to Best Utilize the ever-increasing (\uparrow) Wealth of Resources?

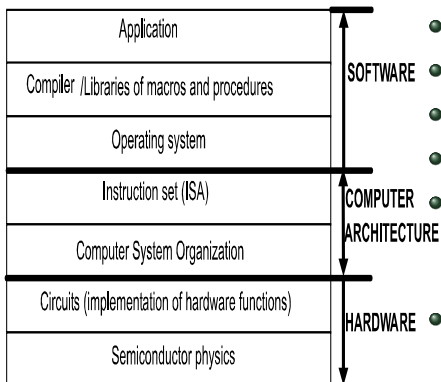


Brief History

- ICPP, ISCA 1980/90s: parallel architectures popular topic.
Demise of SingleCPU System: inevitable & fast approaching.
- **Whatever happened? Mid90 Killer-Micro:**
 - The rapid increase (\uparrow) transistor density \Rightarrow
 - path of least resistance: ever-increasing speed of SingleCPU
 - Complex Out-Of-Order (OoO) processors: 100s instructions/cycle.
 - Commercial arena: multiprocessors just an uniprocessor extension.
- **What Changed?** Multiprocessors Trend: Academia & Industry:
 - power complexity
 - **Memory WALL:** \uparrow performance gap between processor & memory
- All Future Architectures adopt some form of massive parallelism!



What used to Be Computer Architecture?



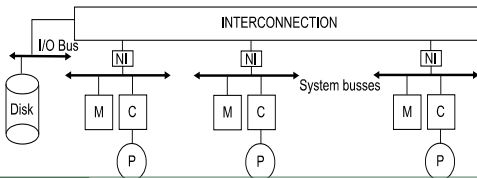
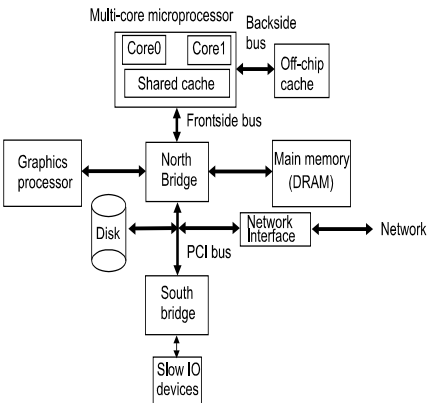
- Multi-Layered: focus competences.
- Each layer uses the ones below it.
- **Application** C++/Java/SML/Haskell.
- **Compiler**: machine code + OpSys calls.
- **OpSys**: extends hardware funct & orchestrates resource sharing among multiple users.
- **ISA & Computer Organization**: software-hardware interface.

Old Definition ISA – critical role in the success of computer industry:

- Early on, used to be the hallmark of system design ⇒
- Non-Portable Software & No Compiler at that time
- 1960s IBM System360 ISA guarantees backward compatibility ⇒
- Strategy endured test of time & Behemoth company today



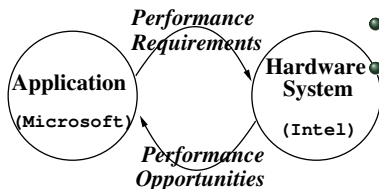
What Is Computer Architecture Today?



- **Much Broader Definition:** how best to build the computer (includes ISA, but focus changed on organization).
- North Bridge: sys bus connects cores to memory & IO devs.
- PCI bus: IO bus connecting North Bridge to high-speed IO to disk, network & slower dev
- ← Generic Multiprocessor with Distributed Memory
- **Parallel Sys Main Components**
(1) Processor, (2) Memory Hierarchy and (3) Interconnection



Software-Hardware Synergy & Biggest Challenge



- Computer Architect Role:
 - design trade-offs across HW/SW interf to meet functional & performance requirements within cost constraints.

- software: flexible approach to simplifying hardware, e.g., TLB exceptions, FP ops,
- hardware faster \Rightarrow world of tradeoff in between
- guided by the common case \Rightarrow trends are important.
- Important Juncture:
 - Higher performance requires Parallel Hardware
 - Biggest Challenge: develop efficient Massively-Parallel Software!



Course Organization

W	HARDWARE		SOFTWARE	LAB/CUDA
1	Trends Vector Machine	←	List HOM (Map-Reduce)	Intro & Simple Map Programming
2	In Order Processor	→ ←	VLIW Instr Scheduling	Scan & Reduce
3	Cache Coherence		Reasoning About Parallelism	Sparse Vect Matrix Mult
4	Interconnection Networks		Case Studies & Optimizations	Transpose & Matrix Matrix Mult
5	Memory Consistency		Optimising Locality	Sorting & Profiling & Mem Optimizations
6	OoO, Spec Processor		Thread-Level Speculation	Project Work

Three narrative threads: the path to complex & good design:

- Design Space tradeoffs, constraints, common case, trends.
- Reasoning: from simple to complex, Applying Concepts.



1 Introduction²

- Brief History
- Computer Architecture Definition
- Course Organization

2 Trends of Critical Components of a Parallel System

- Processor
- Memory
- Interconnect

3 Technological Challenges/Constraints

- Power
- Reliability
- Wire Delays
- Design Complexity
- CMOS Endpoint

4 How to Reason about Parallelism?

- Amdahl's Law
- List Homomorphism
- List-Homomorphism \equiv Map-Reduce



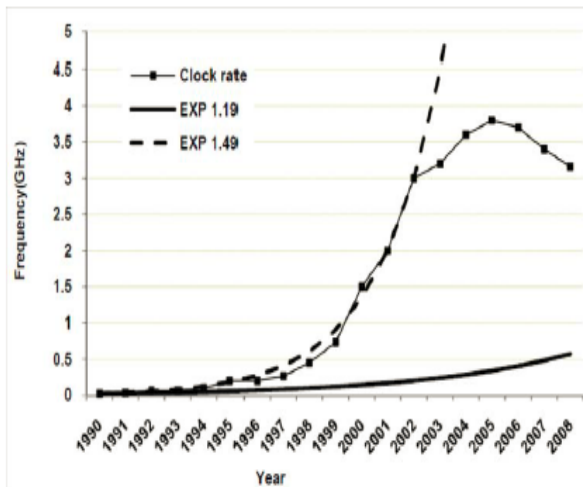
Abstractions

- A **program** is to a **process/thread** what a recipe is for cooking.
- **Processor (core)**: hardware entity capable of sequencing & executing thread's instructions.
- **MT Cores** multiple threads, each running in its thread context.
- **Multiprocessor**: set of processors connected to execute a workload
 - mass produced, off-the-shelf, each several cores & levels of cache
 - trend towards migrating system functions on the chip:
memory controllers, external cache directories, network interface



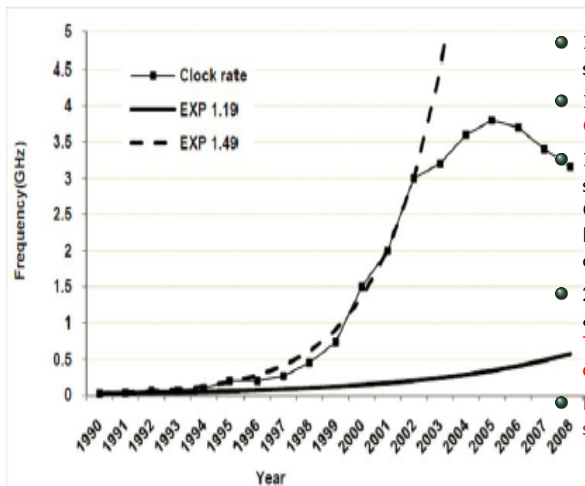
Processor: Clock Frequency/Rate

Historically the clock rate (at which instr are executed) has increased exponentially (1990-2004).



Processor: Clock Frequency/Rate

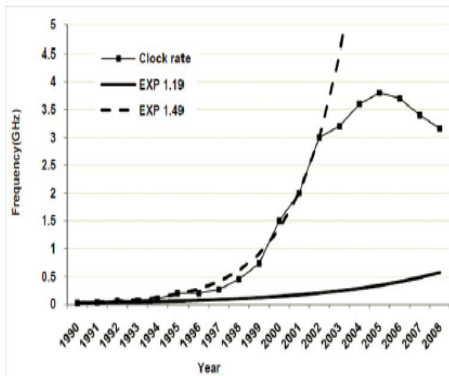
Historically the clock rate (at which instr are executed) has increased exponentially (1990-2004).



- 1.19× per-year due technology scaling (same hwd on new techn).
- 1990-2002: doubled every 21 months
Curve 1.49×: 30GHz in 2008!
- 1.49 × — 1.19×: very-deep (10-20 stages) pipelines! ILP via speculative OoO: register renaming, reorder bufs, branch prediction, lockup-free caches, memory disambiguation, etc.
- 2004: Intel cancels Pentium4 @4Ghz & Switched Track to Multi-Cores ⇒ **Tectonic Shift away from muscled deeply-pipelined uniprocessor.**
- Peaked in 2005, but mostly stalled since 2002!



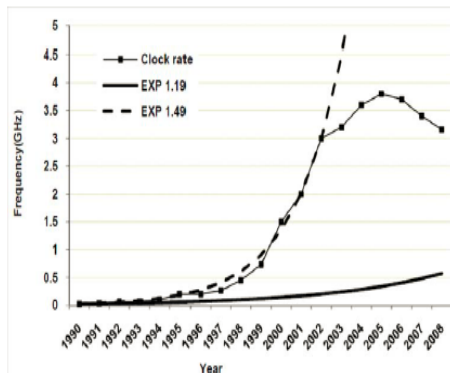
Closer Look at Clock Rate



- Technology (process shrinkage): every generation transistors' switching speed increases 41%.
- Pipeline Depth: more stages \Rightarrow less complex \Rightarrow less gates/stage
 - # of gates delays dropped by 25% every process generation.
- Improved Circuit Design



Closer Look at Clock Rate



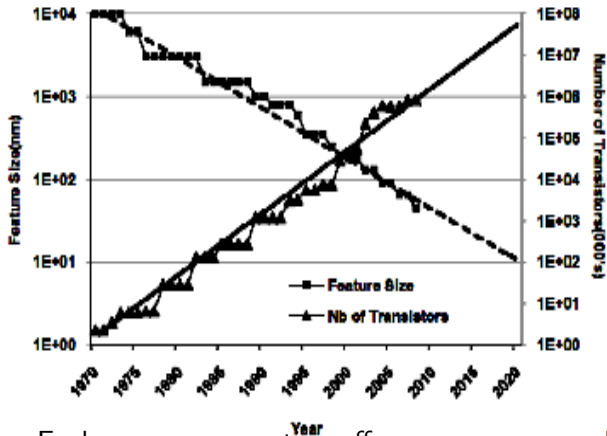
- **Technology (process shrinkage):** every generation transistors' switching speed increases 41%.
- **Pipeline Depth:** more stages \Rightarrow less complex \Rightarrow less gates/stage
 - # of gates delays dropped by 25% every process generation.
- **Improved Circuit Design**

Clock Rate Increase is Not Sustainable:

- **Deeper pipelines:** difficult to build useful stages with < 10 gates
- **Wire delays:** wire-transm speed \uparrow much slower than switching,
- Circuits clocked at **higher rates consume more power!**



Processor: Feature Size & Number of Transistors



- new process generation every 2 years
- feature size reduced 30% every generation
- # of transistors doubles every 2 years (Moore's law).
1 Billion in 2008.

Each process generation offers new resources. How best to use the > 100 billion transistors? Large-Scale CMPs (100s-1000s cores):

- more cache, better memory-system design
- fetch and decode multiple instr per clock
- running multiple threads per core and on multiple cores

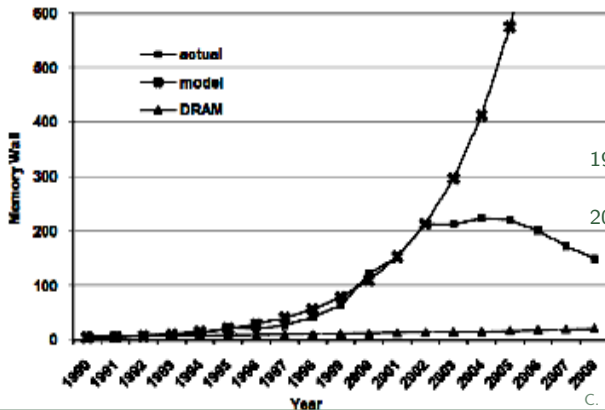
Memory Systems

- **(Main) Memory Wall:** growing gap between processor and memory speed. Processor cannot execute faster than memory system can deliver data and instructions!
- Want Big, Fast & Cheap Memory System
 - access time increases with memory size as it is dominated by wire delays \Rightarrow this will not change in future technologies
 - multi-level hierarchies (relies on principle of locality)
 - efficient management is KEY, e.g., cache coherence.
 - Cost and Size memories in a basic PC in 2008:

Memory	Size	Marginal Cost	Cost Per MB	Access Time
L2 Cache	1MB	\$20/MB	\$20	5nsec
Main Memory	1 GB	\$50/GB	5c	200 nsec
Disk	500GB	\$100/500GB	0.02c	5 msec

Memory Wall? Which Memory Wall??

- DRAM density increases $4\times$ every 3 years, BUT
- DRAM speed \uparrow only with 7% per year! (processor speed by 50%)
- Perception was that Memory Wall will last forever!
- Memory Wall Stopped Growing around 2002.
- Multi/Many-Cores \Rightarrow shifted from Latency to Bandwidth WALL

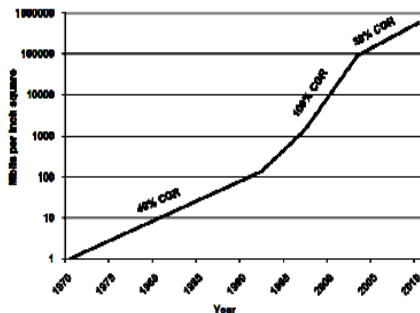
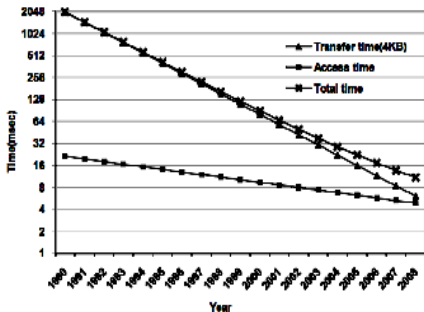


- $\text{MemoryWall} = \frac{\text{mem_cycle}}{\text{proc_cycle}}$
- 1990 MemoryWall = 4 (25MHz, 150ns)
- 2002 exponential growth MemoryWall = 200
- Stalled since then.
- If trend continues: 1 Terabit Main Memory by 2021.



Disk Memory

- Historically disk performance improved by 40% per year
- $\text{DiskTime} = \text{AccessTime} + \text{TransferTime}$ ($\text{AccessTime} = \text{Seek} + \text{Latency}$)
- Historically, transfer time have dominated, but
- Today: transfer and access time are of the same **msecs** order
- Future, Access Time will dominate, but proc-disk gap still large



Seek Time: head to reach right track, latency: time to reach the first record on track, both depend on rotation speed & independent on block size.



Interconnection Networks

Present at many layers:

- **On-Chip Interconnects:** forward values between pipeline stages, AND between execution units AND connect cores to shared cache banks.
- **System Interconnects:** connect processors (CMPs) to memory and IO
- **I/O Interconnects,** usually bus e.g., PCI, connect various devices to the System Bus
- **Inter-Systems Interconnects:** connect separate systems (chassis or boxes) & include
 - **SANs:** connect systems at very short distance
 - LANs, WANs (not interesting for us).
- Internet: global world-wide interconnect (not interesting for us).



1 Introduction²

- Brief History
- Computer Architecture Definition
- Course Organization

2 Trends of Critical Components of a Parallel System

- Processor
- Memory
- Interconnect

3 Technological Challenges/Constraints

- Power
- Reliability
- Wire Delays
- Design Complexity
- CMOS Endpoint

4 How to Reason about Parallelism?

- Amdahl's Law
- List Homomorphism
- List-Homomorphism \equiv Map-Reduce



Technological Constraints

- In the Past: tradeoff between cost (area) and time (performance).
- Today: design is challenged by several technological limits
 - Major new constraint is Power
 - wire delays
 - reliability
 - complexity of design
- It seems that parallelism addresses well all these constraints.



Power

- Total Power = Dynamic + Static (Leakage)

$P_{dynamic} = \alpha CV^2 f$ consumed by a gate when it switches state

$P_{static} = VI_{sub} \sim Ve^{-kV_T/T}$ (caches)

- Dynamic power favors parallel processing over higher clock rate
 - $P_{dynamic} \sim F^3$ mostly dissipated in processor



Power

- Total Power = Dynamic + Static (Leakage)

$P_{dynamic} = \alpha CV^2 f$ consumed by a gate when it switches state

$P_{static} = VI_{sub} \sim Ve^{-kV_T/T}$ (caches)

- Dynamic power favors parallel processing over higher clock rate
 - $P_{dynamic} \sim F^3$ mostly dissipated in processor
 - increase clock freq $4\times \Rightarrow$



Power

- **Total Power = Dynamic + Static (Leakage)**

$P_{dynamic} = \alpha CV^2 f$ consumed by a gate when it switches state

$P_{static} = VI_{sub} \sim Ve^{-kV_T/T}$ (caches)

- Dynamic power favors parallel processing over higher clock rate
 - $P_{dynamic} \sim F^3$ mostly dissipated in processor
 - increase clock freq $4\times \Rightarrow 4\times$ speedup @ $64\times$ dynamic power!
 - replicate a uniprocessor $4\times \Rightarrow 4\times$ speedup @ $4\times$ power
- Static Power: dissipated in all circuits, at all time, nomatter of frequency and whether it switches or not.
 - negligible 15 years ago, but as feature size decreased so did the threshold voltage V_T every generation
 - **Recently overtook dynamic power as major source of dissipation!**
- **Power/Energy are Critical Problems**
e.g., costly & many battery operated devices.



Reliability

- **Transient Failures (Soft Errors):**
 - Charge stored in a transistor $Q = C V$
 - Supply voltage V decreases every generation (consequence of features-size shrinking)
 - As Q decreases, it is easier to flip bits
 - Corruption Sources: cosmic rays, alpha particles radiating from the packaging material, electrical noise
 - Device operational but values have been corrupted
 - DRAM/SRAM error detection and correction capability
- **Intermittent/Temporary Failures:**
 - last longer, should try to continue execution
 - aging or temporary environmental variation, e.g., temperature
- **Permanent Failures:** device will never function again, must be isolated & replaced by spare
- **Chip Mutiprocessors: promote better reliability**
 - using threads for redundant execution,
 - faulty cores can be disabled \Rightarrow natural failsafe degradation



Wire Delays

- Miniaturization \Rightarrow transistors switch faster, but the propagation delay of signals on wire does not scale as well.
- Wire Delay Propagation $\sim RC$. $R \sim L/CS_{area}$. Miniaturization \Rightarrow cross-section area keeps shrinking each generation, annuls the benefit of length shrinking.
- Wires can be pipelined like logic.
- Deeper pipelines are better because communication limited to only few stages.
- Impact of wire delays also favors multiprocessors, since communication traffic is hierarchical:
 - most communication is local
 - inter-core communication is occasional



Design Complexity

- Design Verification has become the dominant cost of chip development today, major design constraint.
- Chip density increases much faster than the productivity of verification engineers (new tools and speed of systems):
 - register-transfer language level, i.e., logic is correct
 - core level, i.e., correctness of forwarding, memory disambiguation,
 - multi-core level, e.g., cache coherence, memory consistency.
- Vast majority of chip resources dedicated to storage also due to verification complexity:
 - trivial to increase the size of: caches, store buffers, load/store/fetch queues, reorder buffers, directory for cache coherence, etc.
- Design Trend Favors Multiprocessors: easier to replicate the same structure multiple times than to design a large, complex one.



CMOS (Endpoint) Meets Quantum Physics

- CMOS is rapidly reaching the limits of miniaturization,
- Feature size: half pitch distance (half the distance between two metal wires). Gate length $\sim 1/2$ feature size.
- If present trends continues feature size $< 10nm$ by 2020
- Radius of atom: $0.1 \sim 0.2nm \Rightarrow$ gate length quickly reaches atomic distances that are governed by quantum physics, i.e., binary logics replaced with probabilistic states.
- Not clear what will follow (?)



1 Introduction²

- Brief History
- Computer Architecture Definition
- Course Organization

2 Trends of Critical Components of a Parallel System

- Processor
- Memory
- Interconnect

3 Technological Challenges/Constraints

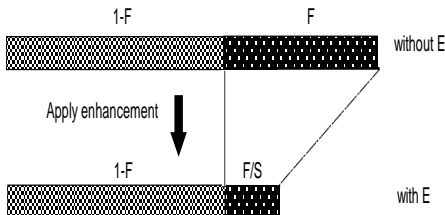
- Power
- Reliability
- Wire Delays
- Design Complexity
- CMOS Endpoint

4 How to Reason about Parallelism?

- Amdahl's Law
- List Homomorphism
- List-Homomorphism \equiv Map-Reduce



Amdahl's Law



Enhancement accelerates a fraction F of the task by a factor S :

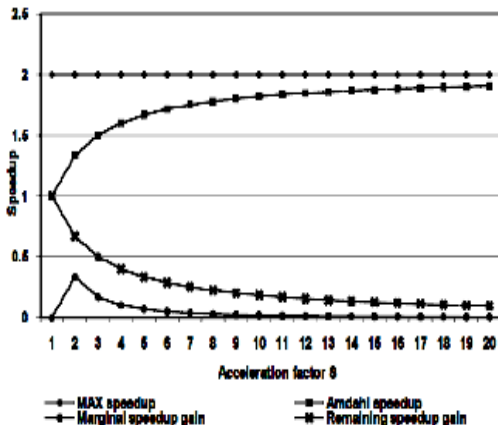
$$T_{\text{exe}}(\text{with}E) = T_{\text{exe}}(\text{without}E) \times \left[(1 - F) + \frac{F}{S} \right]$$

$$\text{Speedup}(E) = \frac{T_{\text{exe}}(\text{without}E)}{T_{\text{exe}}(\text{with}E)} = \frac{1}{(1-F) + \frac{F}{S}}$$



Amdahl's Law

- 1 Improvement is limited by the $1 - F$ part of the execution that cannot be optimized: $Speedup(E) < \frac{1}{1-F}$
- 2 Optimize the common case & execute the rare case in software.
- 3 Low of diminishing returns

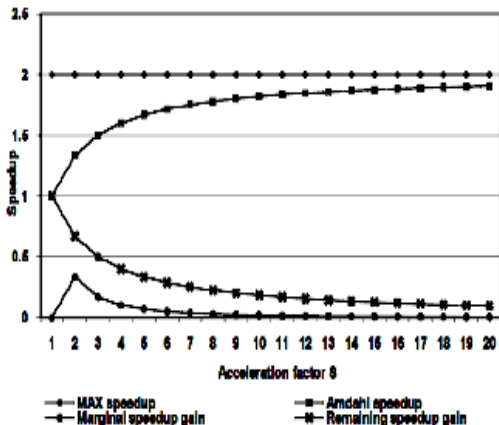


$F=0.5$



Amdahl's Law

- 1 Improvement is limited by the $1 - F$ part of the execution that cannot be optimized: $Speedup(E) < \frac{1}{1-F}$
- 2 Optimize the common case & execute the rare case in software.
- 3 Low of diminishing returns



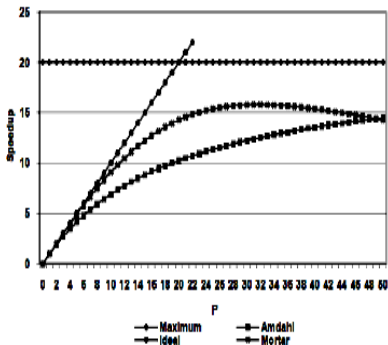
$F=0.5$

- every increment of S consumes new resources and is less rewarding:
- $S = 2 \Rightarrow 33\%$ speedup,
- $S = 5 \Rightarrow 6.67\%$ speedup.



Amdahl's Law: Parallel Speedup

$$S_P = \frac{T_1}{T_P} = \frac{P}{F + P(1-F)} < \frac{1}{1-F}$$



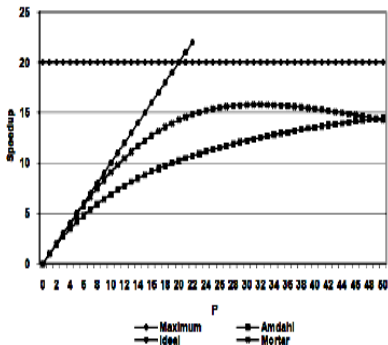
F=0.95

- Typically: speedup is sublinear, e.g., due to inter-thread communic.
- Sometimes superlinear speedup due to cache effects.
- Unforgiving Law: even if 99% is parallelized, $S_{\infty} < 100$.



Amdahl's Law: Parallel Speedup

$$S_P = \frac{T_1}{T_P} = \frac{P}{F + P(1-F)} < \frac{1}{1-F}$$



F=0.95

- Typically: speedup is sublinear, e.g., due to inter-thread communic.
- Sometimes superlinear speedup due to cache effects.
- Unforgiving Law: even if 99% is parallelized, $S_{\infty} < 100$.

Hardware Trend is to ever increase the number of cores.

Amdahl's Law: reason about parallelism asymptotically (∞ # cores), i.e., systematically exploit all levels of application's parallelism.



1 Introduction²

- Brief History
- Computer Architecture Definition
- Course Organization

2 Trends of Critical Components of a Parallel System

- Processor
- Memory
- Interconnect

3 Technological Challenges/Constraints

- Power
- Reliability
- Wire Delays
- Design Complexity
- CMOS Endpoint

4 How to Reason about Parallelism?

- Amdahl's Law
- List Homomorphism
- List-Homomorphism \equiv Map-Reduce



Math Preliminaries: Monoid & Homomorphism

Definition (Monoid)

Assume set S and $\odot : S \times S \rightarrow S$. (S, \odot) is called a *Monoid* if it satisfies the following two axioms:

- (1) *Associativity*: $\forall x, y, z \in S$ we have $(x \odot y) \odot z \equiv x \odot (y \odot z)$ and
- (2) *Identity Element*: $\exists e \in S$ such that $\forall a \in S$, $e \odot a \equiv a \odot e \equiv a$.

$((S, \odot)$ is called a *group* if it also satisfies that any element is invertible, i.e., $\forall a, \exists a^{-1}$ such that $a \odot a^{-1} \equiv a^{-1} \odot a \equiv e$.)

E.g., $(\mathbb{N}, +)$, (\mathbb{Z}, \times) , $(\mathbb{L}_T, ++)$, where

\mathbb{L}_T denotes lists of elements of type T , and $++$ list concatenation.



Math Preliminaries: Monoid & Homomorphism

Definition (Monoid)

Assume set S and $\odot : S \times S \rightarrow S$. (S, \odot) is called a *Monoid* if it satisfies the following two axioms:

- (1) *Associativity*: $\forall x, y, z \in S$ we have $(x \odot y) \odot z \equiv x \odot (y \odot z)$ and
- (2) *Identity Element*: $\exists e \in S$ such that $\forall a \in S, e \odot a \equiv a \odot e \equiv a$.

$((S, \odot)$ is called a *group* if it also satisfies that any element is invertible, i.e., $\forall a, \exists a^{-1}$ such that $a \odot a^{-1} \equiv a^{-1} \odot a \equiv e$.)

E.g., $(\mathbb{N}, +)$, (\mathbb{Z}, \times) , $(\mathbb{L}_T, ++)$, where

\mathbb{L}_T denotes lists of elements of type T , and $++$ list concatenation.

Definition (Monoid Homomorphism)

A *monoid homomorphism* from monoid (S, \oplus) to monoid (T, \odot) is a function $h : S \rightarrow T$ such that $\forall u, v \in S, h(u \oplus v) \equiv h(u) \odot h(v)$.



List Homomorphism (LH)

- Finite lists with concatenation, denoted $++$, forms a Monoid:
 - $[1,2,3,4] ++ [5,6,7,8] = [1,2,3,4,5,6,7,8]$,
 - neutral elem e_{++} is the empty list $[]$, i.e., $[1,2] ++ [] = [1,2]$,
 - later we will look at them as **vectors** rather than **linked lists**.



List Homomorphism (LH)

- Finite lists with concatenation, denoted $++$, forms a Monoid:
 - $[1,2,3,4] ++ [5,6,7,8] = [1,2,3,4,5,6,7,8]$,
 - neutral elem e_{++} is the empty list $[]$, i.e., $[1,2] ++ [] = [1,2]$,
 - later we will look at them as **vectors** rather than **linked lists**.

Definition (List Homomorphism)

$h :: [T_1] \rightarrow T_2$ over finite lists is a list homomorphism if there exists a (n associative) binary operator $\odot :: T_2 \rightarrow T_2 \rightarrow T_2$, such that:

$$h(x ++ y) = (h\ x) \odot (h\ y)$$

We denote $h = \text{hom}(\odot) f\ e$, where $f(x) = h([x])$ and $e = h([])$.

- If the program is well behaved, i.e., same result for any $x ++ y$ splitting of the input list, then (T_2, \odot) is necessarily a Monoid.
- Similar to a divide and conquer algorithm, where $h(x)$ and $h(y)$ can be computed in parallel and results can be merged.



List Homomorphism (LH) Examples

Definition (List Homomorphism)

$h :: [T_1] \rightarrow T_2$ over finite lists is a list homomorphism if there exists an associative binary operator $\odot :: T_2 \rightarrow T_2 \rightarrow T_2$, such that:

$$h (x++y) = (h x) \odot (h y)$$

We denote $h = \text{hom } (\odot) f e$, where $f x = h [x]$ and $e = h []$.

Examples of List Homomorphisms, $\text{id } x = x$ is the identity function

```
-- len = hom (+) one 0, one x = 1
len :: [T] -> Integer
len []      = 0
len [x]     = 1
len (x++y) = (len x) + (len y)
```

```
-- flatten = hom (++) id []
flatten :: [[T]] -> [T]
flatten []      = []
flatten [x]     = x
flatten (x++y) = (flatten x) ++
                  (flatten y)
```

```
-- maxList = hom (max) id -∞
maxList :: [Int] -> Int
maxList []      = -∞
maxList [x]     = x
maxList (x++y) = (maxList x) 'max'
                  (maxList y)

-- Assume p :: T -> Bool given,
-- all_p = hom (&&) p True
all_p :: [T] -> Bool
all_p []      = ???
all_p [x]     = ???
all_p (x++y) = ???
```

Exercise: Implement the above LHs in Haskell

Implementation Sample for all_p :

```
import System.Environment -- access to arguments etc.
-- helper to simulate (x++y) pattern
split :: [a] -> ([a], [a])
split [] = ([], [])
split [x] = ([], [x])
split xs = let mid = (length xs) `div` 2 in (take mid xs, drop mid xs)

--  $\text{all}_p \equiv \text{alln } p = \text{hom } (\&\&) \text{ p True}$ 
alln :: (a -> Bool) -> [a] -> Bool
alln p [] = True
alln p [x] = p x
alln p xs = let (x, y) = split xs in (alln p x) && (alln p y)

----- Compile with: ghc -O2 -o test LHegHaskell.hs -----
main :: IO()
main = do args <- getArgs
        let inp = if null args then [0,2,4,6] else read (head args)
            p x = x `mod` 2 == 0
            res = alln p inp
        putStrLn ("Computed: " ++ show res)
```



1 Introduction²

- Brief History
- Computer Architecture Definition
- Course Organization

2 Trends of Critical Components of a Parallel System

- Processor
- Memory
- Interconnect

3 Technological Challenges/Constraints

- Power
- Reliability
- Wire Delays
- Design Complexity
- CMOS Endpoint

4 How to Reason about Parallelism?

- Amdahl's Law
- List Homomorphism
- List-Homomorphism \equiv Map-Reduce



Basic Blocks of Parallel Programming: Map

map :: $((\alpha \rightarrow \beta), [\alpha]) \rightarrow [\beta]$ has *inherently parallel semantics*.

$$\begin{array}{rcl}
 x & = & \text{map}(f, \{ a_1, a_2, \dots, a_n \}) \\
 & & \downarrow \quad \downarrow \quad \downarrow \\
 x & \equiv & \{ f(a_1), f(a_2), \dots, f(a_n) \}
 \end{array}$$

Map Fusion: (higher-order transformation)

$$\begin{array}{rcl}
 a & = & \{ a_1, a_2, \dots, a_n \} \\
 x & = & \text{map}(f, a) \\
 y & = & \text{map}(g, x) \quad \equiv \quad y = \text{map}(g \circ f, a) \\
 & & \downarrow \\
 y & \equiv & \{ g(f(a_1)), g(f(a_2)), \dots, g(f(a_n)) \} \quad \equiv \quad \{ g(f(a_1)), g(f(a_2)), \dots, g(f(a_n)) \}
 \end{array}$$

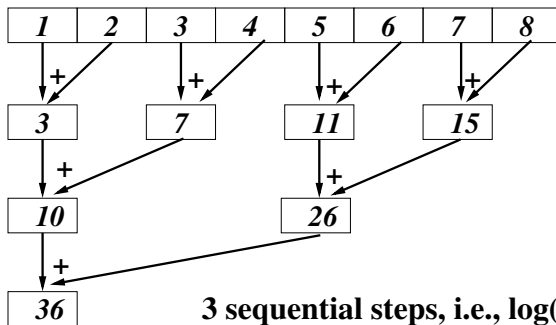


Basic Blocks of Parallel Programming: Reduce

reduce :: $((\alpha \rightarrow \alpha \rightarrow \alpha), \alpha, [\alpha]) \rightarrow \alpha$

reduce(\odot , e , $\{a_1, a_2, \dots, a_n\}$) $\equiv e \odot a_1 \odot a_2 \odot \dots \odot a_n$

where \odot is an associative binary operator.



Build programs by combining **map**, **reduce** and other such operators.



Map, Reduce, and Scan Types and Semantics

- \bullet $\text{map} :: (\alpha \rightarrow \beta) \rightarrow [\alpha] \rightarrow [\beta]$
 $\text{map } f \ [x_1, \dots, x_n] = [f(x_1), \dots, f(x_n)],$
 i.e., $x_i :: \alpha, \forall i$, and $f :: \alpha \rightarrow \beta$.
- \bullet $\text{reduce} :: (\alpha \rightarrow \alpha \rightarrow \alpha) \rightarrow \alpha \rightarrow [\alpha] \rightarrow \alpha$
 $\text{reduce } \odot \ e \ [x_1, x_2, \dots, x_n] = e \odot x_1 \odot x_2 \odot \dots \odot x_n,$
 i.e., $e :: \alpha, x_i :: \alpha, \forall i$, and $\odot :: \alpha \rightarrow \alpha \rightarrow \alpha$.
- \bullet $\text{scan}^{\text{exc}} :: (\alpha \rightarrow \alpha \rightarrow \alpha) \rightarrow \alpha \rightarrow [\alpha] \rightarrow [\alpha]$
 $\text{scan}^{\text{exc}} \odot \ e \ [x_1, \dots, x_n] = [e, e \odot x_1, \dots, e \odot x_1 \odot \dots \odot x_{n-1}]$
 i.e., $e :: \alpha, x_i :: \alpha, \forall i$, and $\odot :: \alpha \rightarrow \alpha \rightarrow \alpha$.
- \bullet $\text{scan}^{\text{inc}} :: (\alpha \rightarrow \alpha \rightarrow \alpha) \rightarrow \alpha \rightarrow [\alpha] \rightarrow [\alpha]$
 $\text{scan}^{\text{inc}} \odot \ e \ [x_1, \dots, x_n] = [e \odot x_1, \dots, e \odot x_1 \odot \dots \odot x_n]$
 i.e., $e :: \alpha, x_i :: \alpha, \forall i$, and $\odot :: \alpha \rightarrow \alpha \rightarrow \alpha$.



1st List-Homomorphism Theorem [Meertens]

Theorem (1st List Homomorphism Theorem (Meertens))

Any homomorphism $h = \text{hom } (\odot) f e$ can be written as the functional composition of a reduce and a map:

$$h = \text{hom } (\odot) f e_{\odot} = (\text{reduce } (\odot) e_{\odot}) . (\text{map } f)$$

Conversely, each such composition is a homomorphism.

Apply Theorem to Re-Write LH to Map-Reduce!

```
-- len = hom (+) one 0, one x = 1
len :: [T] -> Integer
len []      = 0
len [x]     = 1
len (x++y) = (len x) + (len y)
```

```
-- flatten = hom (++) id []
flatten :: [[T]] -> [T]
flatten []      = []
flatten [x]     = x
flatten (x++y) = (flatten x) ++
                  (flatten y)
```

```
-- maxList = hom (max) id -∞
maxList :: [Int] -> Int
maxList []      = -∞
maxList [x]     = x
maxList (x++y) = (maxList x) 'max'
                  (maxList y)

-- Assume p :: T -> Bool given,
-- all_p = hom (&&) p True
all_p :: [T] -> Bool
all_p []      = ???
all_p [x]     = ???
all_p (x++y) = ???
```



1st List-Homomorphism Theorem [Meertens]

Theorem (1st List Homomorphism Theorem (Meertens))

Any homomorphism $h = \text{hom } (\odot) f e$ can be written as the functional composition of a reduce and a map:

$$h = \text{hom } (\odot) f e_{\odot} = (\text{reduce } (\odot) e_{\odot}) . (\text{map } f)$$

Conversely, each such composition is a homomorphism.

Theorem tells how to parallelize LHs based on map-reduce skeletons.

Map-Reduce Definition for the Discussed LH Examples

```
-- len = hom (+) one 0, one x = 1
len = (reduce (+) 0) . (map one)

-- sum = hom (+) id 0, id x = x
sum = (reduce (+) 0) . (map id)

-- flatten = hom (++) id [],
flatten = (reduce (++) []) . (map id)
```

```
-- maxList = hom (max) id -∞
maxList = (reduce (max) (-∞)) .
          (map id)

-- allp = hom (&&) p True
allp = (reduce (&&) True) .
        (map p)
```


List Homomorphism Invariants

Theorem (List-Homomorphism Promotions)

Given unary functions f , g and an associative binary operator \odot then:

1. $(\text{map } f) . (\text{map } g) \equiv \text{map } (f . g)$
2. $(\text{map } f) . (\text{reduce } (++) []) \equiv (\text{reduce } (++) []) . (\text{map } (\text{map } f))$
3. $(\text{reduce } \odot e_{\odot}) . (\text{reduce } (++) []) \equiv$
 $(\text{reduce } \odot e_{\odot}) . (\text{map } (\text{reduce } \odot e_{\odot}))$

- 2. 3. \Rightarrow **code generation**: list is segmented, segments are distributed on different processors, computation proceeds locally on each processor, and the local results are reduced.
- 2. 3. \Leftarrow **flattening optimization**: uncovers more parallelism
- e.g., $\text{map } f [1..4] = (\text{map } f) . (\text{red } ++) [[1,2],[3,4]] =^{\text{prom2}} ?$



Optimizing Map-Reduce Computation (**Exercise**)

Theorem (Optimized Map Reduce)

Assume $\text{distr}_p :: [\alpha] \rightarrow [[\alpha]]$ distributes a list into p sublists, each containing about the same number of elements. Denoting $\text{redomap} \odot f e_{\odot} \equiv (\text{reduce} \odot e_{\odot}) . (\text{map } f)$, the equality holds:

$$\text{redomap} \odot f e_{\odot} \equiv (\text{reduce} \odot e_{\odot}) . (\text{map } (\text{redomap} \odot f e_{\odot})) . \text{distr}_p$$

- Prove it using the promotion Lemmas before!
- Hint: $(\text{reduce } (++) []) . \text{distr}_p \equiv \text{id}$, hence
- $\text{redomap} \odot f e_{\odot} \equiv (\text{reduce} \odot e_{\odot}) . (\text{map } f) . (\text{reduce } (++) []) . \text{distr}_p$



Are All List Homomorphism Efficient?

If the combine operator involves concatenation then does map-reduce provides efficient parallelization?

Merge Sort

```
-- merge two sorted lists
merge :: Ord T => [T] -> [T] -> [T]
merge [] y = y
merge x [] = x
merge (x:xs) (y:ys) =
  if ( x <= y )
  then x : merge xs (y:ys)
  else y : merge (x:xs) ys
```

```
-- mSort = hom merge [.] []
-- [.] x = [x]
mSort :: Ord T => [T] -> [T]
mSort [] = []
mSort [x] = [x]
mSort (x++y) = (mSort x) 'merge'
               (mSort y)
```

In the naive merged sort, the merge reduction operator traverses sequentially the whole list, hence this map-reduce does not give efficient parallelization!



Conclusion

What have we talked about today?

- Hardware Parallelism:
the only way of scalably increasing the compute power.
 - demonstrated by hardware trends:
 - power, reliability, wire delays, design complexity.
- Big Challenge: having parallel commodity software.
- List-Homomorphism:
a way of reasoning about parallelism and
of building inherently parallel programs.

