



Faculty of Science



Loop Parallelism I

Cosmin E. Oancea and Troels Henriksen
[cosmin.oancea,athas]@diku.dk

Department of Computer Science (DIKU)
University of Copenhagen

September 2015 PMPH Lecture Notes



Course Organization

W	HARDWARE		SOFTWARE	LAB/CUDA
1	Trends Vector Machine	←	List HOM (Map-Reduce)	Intro & Simple Map Programming
2	In Order Processor	→ ←	VLIW Instr Scheduling	Scan & Reduce
3	Cache Coherence		Loop Parallelism I	Sparse Vect Matrix Mult
4	Interconnection Networks		Case Studies & Optimizations	Transpose & Matrix Matrix Mult
5	Memory Consistency		Optimising Locality	Sorting & Profiling & Mem Optimizations
6	OoO, Spec Processor		Thread-Level Speculation	Project Work

Three narrative threads: the path to complex & good design:

- **Design Space** tradeoffs, constraints, common case, trends.
- **Reasoning**: from simple to complex, **Applying Concepts**.



Motivation

- + So far we reasoned about how to parallelize a known algorithm
- + using a clean, functional approach, e.g., flattening,
- + which provides work and depth guarantees,
 - but does **NOT** account for locality of reference.

Why do we have to look at imperative loops?



Motivation

- + So far we reasoned about how to parallelize a known algorithm
- + using a clean, functional approach, e.g., flattening,
- + which provides work and depth guarantees,
 - but does **NOT** account for locality of reference.

Why do we have to look at imperative loops?

- A lot of legacy sequential imperative code, C++/Java/Fortran.
- Need to parallelize the implementation of unknown algorithm,
- Need to optimize parallelism, e.g., locality of reference requires subscript analysis.



- 1 Direction-Vector Analysis
- 2 Block Tiling: Matrix Multiplication Case Study
- 3 Coalesced Accesses: Matrix Transposition Case Study
- 4 Non-Trivial Synchronization: Histogram Case Study
- 5 Known Recurrences: Tridiagonal Solver Case Study
- 6 Imperative Context: Summarization of Array Indexes



Problem Statement

Three Loop Examples

DO i = 1, N	DO i = 2, N	DO i = 2, N
DO j = 1, N	DO j = 2, N	DO j = 1, N
A[j,i] = A[j,i] ..	A[j,i] = A[j-1,i-1]...	A[i,j] = A[i-1,j+1]...
ENDDO	B[j,i] = B[j-1,i]...	ENDDO
ENDDO	ENDDO ENDDO	ENDDO

Iterations are ordered *lexicographically*, i.e., in the order they occur in the sequential execution, e.g., $\vec{k} = (i=2, j=4) < \vec{l} = (i=3, j=3)$.

- Which of the three loop nests is amenable to parallelization?
- Loop interchange is one of the most simple and useful code transformations, e.g., used to enhance locality of reference, parallel-loop granularity, and even to “create” parallelism.
- In which loop nest is it safe to interchange the loops?



Definition of a Dependency

Load-Store Classification of Dependencies

True Dependency (RAW)

S1 X = ..

S2 .. = X

Anti Dependency (WAR)

S1 .. = X

S2 X = ..

Output dependency (WAW)

S1 X = ...

S2 X = ...

Th. Loop Dependence: There is a dependence from statement $S1$ to $S2$ in a loop nest *iff* \exists iterations \vec{k}, \vec{l} such that:



Definition of a Dependency

Load-Store Classification of Dependencies

True Dependency (RAW)		Anti Dependency (WAR)		Output dependency (WAW)	
S1	X = ..	S1	.. = X	S1	X = ...
S2	.. = X	S2	X = ..	S2	X = ...

Th. Loop Dependence: There is a dependence from statement $S1$ to $S2$ in a loop nest *iff* \exists iterations \vec{k}, \vec{l} such that:

1. $\vec{k} < \vec{l}$ or $\vec{k} = \vec{l}$ and \exists an execution path from statement $S1$ to statement $S2$ **such that:**
2. $S1$ accesses memory location M on iteration \vec{k} , and
3. $S2$ accesses memory location M on iteration \vec{l} , and
4. one of these accesses is a write.

We say that $S1$ is the source and $S2$ is the sink of the dependence, because $S1$ executes before $S2$ in the sequential program execution. Dependence depicted with an arrow pointing from source to sink.



Definition of a Dependency

Load-Store Classification of Dependencies

True Dependency (RAW)	Anti Dependency (WAR)	Output dependency (WAW)
S1 X = ..	S1 .. = X	S1 X = ...
S2 .. = X	S2 X = ..	S2 X = ...

Th. Loop Dependence: There is a dependence from statement $S1$ to $S2$ in a loop nest *iff* \exists iterations \vec{k}, \vec{l} such that:

1. $\vec{k} < \vec{l}$ or $\vec{k} = \vec{l}$ and \exists an execution path from statement $S1$ to statement $S2$ **such that:**
2. $S1$ accesses memory location M on iteration \vec{k} , and
3. $S2$ accesses memory location M on iteration \vec{l} , and
4. one of these accesses is a write.

We say that $S1$ is the source and $S2$ is the sink of the dependence, because $S1$ executes before $S2$ in the sequential program execution. Dependence depicted with an arrow pointing from source to sink.

We are most interested in cross iteration dependencies, i.e., $\vec{k} < \vec{l}$. Intra iteration dependencies, i.e., $\vec{k} = \vec{l}$ are analysed for ILP.



Loop-Nest Dependencies

Lexicographic ordering, e.g., $\vec{k}=(i=2,j=4) < \vec{l}=(i=3,j=3)$.

Three Loop Examples

```
DO i = 1, N
  DO j = 1, N
    A[j,i] = A[j,i] ..
  ENDDO
ENDDO
```

```
DO i = 2, N
  DO j = 2, N
    A[j,i] = A[j-1,i-1]...
    B[j,i] = B[j-1,i]...
  ENDDO ENDDO
```

```
DO i = 2, N
  DO j = 1, N
    A[i,j] = A[i-1,j+1]...
  ENDDO
ENDDO
```



Loop-Nest Dependencies

Lexicographic ordering, e.g., $\vec{k} = (i=2, j=4) < \vec{l} = (i=3, j=3)$.


Three Loop Examples

```
DO i = 1, N
  DO j = 1, N
    A[j,i] = A[j,i] ..
  ENDDO
ENDDO
```


```
DO i = 2, N
  DO j = 2, N
    A[j,i] = A[j-1,i-1]...
    B[j,i] = B[j-1,i]...
  ENDDO ENDDO
```

```
DO i = 2, N
  DO j = 1, N
    A[i,j] = A[i-1,j+1]...
  ENDDO
ENDDO
```


Leftmost Loop



Middle Loop



Rightmost Loop



How can I summarize this information?



Aggregate Dependencies via Direction Vectors

Write the Direction Vectors for Each Loop:

DO i = 1, N	DO i = 2, N	DO i = 2, N
DO j = 1, N	DO j = 2, N	DO j = 1, N
S1 A[j,i]=A[j,i]..	S1 A[j,i]=A[j-1,i]...	S1 A[i,j]=A[i-1,j+1]...
ENDDO	S2 B[j,i]=B[j-1,i-1]...	ENDDO
ENDDO	ENDDO ENDDO	ENDDO

Dependencies depicted via an edge *from* the stmt that executes first in the loop nest, i.e., *the source*, to the one that executes later, *the sink*.

Def. Dependence Direction: Assume \exists a dependence from $S1$ in iteration \vec{k} to $S2$ in \vec{l} ($\vec{k} \leq \vec{l}$). *Dependence-direction vector* $\vec{D}(\vec{k}, \vec{l})$:

1. $\vec{D}(\vec{k}, \vec{l})_m = "<"$ if $\vec{k}_m < \vec{l}_m$,
2. $\vec{D}(\vec{k}, \vec{l})_m = "="$ if $\vec{k}_m = \vec{l}_m$,
3. $\vec{D}(\vec{k}, \vec{l})_m = ">"$ if $\vec{k}_m > \vec{l}_m$.

If the source is a write and the sink a read then RAW dependency,
if the source is a read then WAR, if both are writes then WAW.



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

<pre> DO i = 1, N DO j = 1, N S1 A[j,i]=A[j,i].. ENDDO ENDDO For S1→S1: (j1,i1)=(j2,i2) i1 = i2 & j1 = j2 Direction matrix: S1→S1: [=,=]</pre>	<pre> DO i = 2, N DO j = 2, N S1 A[j,i]=A[j-1,i]... S2 B[j,i]=B[j-1,i-1]... ENDDO ENDDO S1→S1: (j1,i1)=(j2-1,i2) i1 = i2 & j1 < j2 S2→S2: (j1,i1)=(j2-1,i2-1) i1 < i2 & j1 < j2 Direction matrix: S1→S1: [=,<] S2→S2: [<,<]</pre>	<pre> DO i = 2, N DO j = 1, N S1 A[i,j]=A[i-1,j+1]... ENDDO ENDDO For S1→S1: (i1,j1) = (i1-1,j2+1) i1 < i2 & j1 > j2 Direction matrix: S1→S1: [<,>]</pre>
---	---	--

Th. Parallelism: A loop in a loop nest is parallel *iff* all its directions are either = or there exists an outer loop whose corresp. direction is <.

A direction vector cannot have > as the first non-= symbol, as that would mean that I depend on something in the future.



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

```

DO i = 1, N
  DO j = 1, N
S1    A[j,i]=A[j,i]..
      ENDDO
  ENDDO

```

```

For S1→S1:
  (j1,i1)=(j2,i2)
  i1 = i2 & j1 = j2

```

Direction matrix:

S1→S1: [=,=]

```

DO i = 2, N
  DO j = 2, N
S1    A[j,i]=A[j-1,i]...
S2    B[j,i]=B[j-1,i-1]...
      ENDDO
  ENDDO

```

```

S1→S1: (j1,i1)=(j2-1,i2)
        i1 = i2 & j1 < j2
S2→S2: (j1,i1)=(j2-1,i2-1)
        i1 < i2 & j1 < j2

```

S1→S1: [=,<]

S2→S2: [<,<]

```

DO i = 2, N
  DO j = 1, N
S1    A[i,j]=A[i-1,j+1]...
      ENDDO
  ENDDO

```

```

For S1→S1:
  (i1,j1) = (i1-1,j2+1)
  i1 < i2 & j1 > j2

```

Direction matrix:

S1→S1: [<,>]

Th. Loop Interchange: A column permutation of the loops in a loop nest is legal *iff* permuting the direction matrix in the same way *does NOT* result in a > direction as the leftmost non-= direction in a row.



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

DO i = 1, N	DO i = 2, N	DO i = 2, N
DO j = 1, N	DO j = 2, N	DO j = 1, N
S1 A[j,i]=A[j,i]..	S1 A[j,i]=A[j-1,i]...	S1 A[i,j]=A[i-1,j+1]...
ENDDO	S2 B[j,i]=B[j-1,i-1]...	ENDDO
ENDDO	ENDDO ENDDO	ENDDO
For S1→S1: j1 = j2	For S1→S1: j1 = j2-1	For S1→S1: i1 = i2-1
i1 = i2	i1 = i2	j1 = j2+1
(i2,j2)-(i1,j1)=	(i2,j2)-(i1,j1)=[=,<]	(i2,j2)-(i1,j1)=[<,>]
[=,=]	For S2→S2: j1 = j2-1	
	i1 = i2-1	
	(i2,j2)-(i1,j1)=[<,<]	

Interchange is safe for the first and second nests, but not for the third!

e.g., [=,<] → [<,<] (for the second loop nest)
 [<,<]



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

DO i = 1, N	DO i = 2, N	DO i = 2, N
DO j = 1, N	DO j = 2, N	DO j = 1, N
S1 A[j,i]=A[j,i]..	S1 A[j,i]=A[j-1,i]...	S1 A[i,j]=A[i-1,j+1]...
ENDDO	S2 B[j,i]=B[j-1,i-1]...	ENDDO
ENDDO	ENDDO ENDDO	ENDDO
For S1→S1: j1 = j2	For S1→S1: j1 = j2-1	For S1→S1: i1 = i2-1
i1 = i2	i1 = i2	j1 = j2+1
(i2,j2)-(i1,j1)=	(i2,j2)-(i1,j1)=[=,<]	(i2,j2)-(i1,j1)=[<,>]
[=,=]	For S2→S2: j1 = j2-1	
	i1 = i2-1	
	(i2,j2)-(i1,j1)=[<,<]	

Interchange is safe for the first and second nests, but not for the third!

e.g., [=,<] → [<,<] (for the second loop nest)

[<,<] [<,<]

After interchange, loop j of the second loop nest is parallel.

Corollary: A parallel loop can be always interchanged inwards.



Dependency Graph and Loop Distribution

Def. Dependency Graph: edges from the source of the dependency, i.e., early iteration, to the sink, i.e., later iteration.

Th. Loop Distribution: Statements that are in a dependence cycle remain in one (sequential) loop. The others are distributed to separate loops in graph order; if no cycle then parallel loops.

Vectorization Example: Remember Vector Machines?

```
DO i = 3, N
S1  A[i] = B[i-2] ...
S2  B[i] = B[i-1] ...
ENDDO

For S2→S1: i1 = i2-2, [<]
For S2→S2: i1 = i2-1, [<]
```



Dependency Graph and Loop Distribution


Def. Dependency Graph: edges from the source of the dependency, i.e., early iteration, to the sink, i.e., later iteration.

Th. Loop Distribution: Statements that are in a dependence cycle remain in one (sequential) loop. The others are distributed to separate loops in graph order; if no cycle then parallel loops.

Vectorization Example: Remember Vector Machines?

```
DO i = 3, N
S1  A[i] = B[i-2] ...
S2  B[i] = B[i-1] ...
ENDDO
```

```
For S2→S1: i1 = i2-2, [<]
For S2→S2: i1 = i2-1, [<]
```



```
DO i = 3, N
S2  B[i] = B[i-1] ...
ENDDO
```

```
DOALL i = 3, N
S1  A[i] = B[i-2] ...
ENDDOALL
```

Corollary: It is always legal to distribute a parallel loop; but requires array expansion for local variables or if output dependencies are present.



Loop Distribution May Require Array Expansion

```
float tmp;  
for(i=2; i<N; i++) {  
    tmp = 2*B[i-2];  
    A[i] = tmp;  
    B[i] = tmp+B[i-1]  
}
```

```
float tmp[N];  
for(int i=2; i<N; i++) {  
    tmp[i] = 2*B[i-2];  
    B[i] = tmp[i]+B[i-1];  
}  
  
forall(int i=2; i<N; i++) {  
    A[i] = tmp[i];  
}
```

No matter where `tmp` is declared (inside or outside the loop) it needs to be expanded into an array in order to do loop distribution.

If `tmp` is declared outside the loop then requires **privatization**,



Loop Distribution May Require Array Expansion

```
float tmp;  
for(i=2; i<N; i++) {  
    tmp = 2*B[i-2];  
    A[i] = tmp;  
    B[i] = tmp+B[i-1]  
}
```

```
float tmp[N];  
for(int i=2; i<N; i++) {  
    tmp[i] = 2*B[i-2];  
    B[i] = tmp[i]+B[i-1];  
}  
  
forall(int i=2; i<N; i++) {  
    A[i] = tmp[i];  
}
```

No matter where `tmp` is declared (inside or outside the loop) it needs to be expanded into an array in order to do loop distribution.

If `tmp` is declared outside the loop then requires **privatization**, because it actually causes frequent WAW dependencies. However its value is written before being used within the same iteration. Hence it is semantically equivalent to a locally declared variable, which will remove the output (WAW) dependency.

Distribution requires array expansion of the scalar `tmp`.



False Dependencies (WAR/WAW)

- **Cross-Iteration Anti Dependencies (WAR)** correspond to a read from the array as it was before the loop \Rightarrow can be eliminated by reading from a copy of the array.
- **Cross-Iteration WAW Dependencies (WAW):**
If they correspond to the case in which every **read** from an scalar or array location is covered by a **previous same-iteration write** \Rightarrow can be eliminated **privatization (renaming)**, which semantically moves the declaration of the variable (scalar or array) inside the loop.
- Direction-vectors reasoning is limited to relatively simple loop nests, e.g., difficult to reason about privatization in such a way.



False Dependencies (WAR/WAW)

Anti Dependency (WAR) and Output Dependency (WAW)

```
// OpenMP code:  
// compile with g++ -fopenmp ...  
float tmp = A[1];  
for (i=0; i<N-1; i++)  
S1 A[i] = A[i+1];  
A[N-1] = tmp;  
//S1→S1: i1+1=i2, [<] WAR
```



False Dependencies (WAR/WAW)

Anti Dependency (WAR) and Output Dependency (WAW)

```
// OpenMP code:
// compile with g++ -fopenmp ...
float tmp = A[1];
for (i=0; i<N-1; i++)
S1 A[i] = A[i+1];
  A[N-1] = tmp;
//S1→S1: i1+1=i2, [<] WAR
```

```
// Solution: copy A into A'
// and use A' for the reads!
float Acopy[N];
#pragma omp parallel for
  for(i=0; i<N; i++) {
    Acopy[i] = A[i];
  }
tmp = A[1];
#pragma omp parallel for private(i)
  for (i=0; i<N-1; i++) {
    A[i] = Acopy[i+1];
  }
A[N-1] = tmp;
```

```
int A[M];
for(i=0; i<N; i++){
  for(int j=0, j<M; j++)
    A[j] = (4*i+4*j) % M;
  for(int k=0; k<N; k++)
    X[i,k]=X[i,k-1] * A[A[(2*i+k)%M]];
}
```

// The write to A[j] causes multiple WAWs,



False Dependencies (WAR/WAW)

Anti Dependency (WAR) and Output Dependency (WAW)

```
// OpenMP code:
// compile with g++ -fopenmp ...
float tmp = A[1];
for (i=0; i<N-1; i++)
S1 A[i] = A[i+1];
A[N-1] = tmp;
//S1→S1: i1+1=i2, [<] WAR
```

```
// Solution: copy A into A'
// and use A' for the reads!
float Acopy[N];
#pragma omp parallel for
  for(i=0; i<N; i++) {
    Acopy[i] = A[i];
  }
tmp = A[1];
#pragma omp parallel for private(i)
  for (i=0; i<N-1; i++) {
    A[i] = Acopy[i+1];
  }
A[N-1] = tmp;
```

```
int A[M];
for(i=0; i<N; i++){
  for(int j=0, j<M; j++)
    A[j] = (4*i+4*j) % M;
  for(int k=0; k<N; k++)
    X[i,k]=X[i,k-1] * A[A[(2*i+k)%M]];
}
```

```
// The write to A[j] causes multiple WAWs,
// but A is fully written in the inner loop
#pragma omp parallel{
  int A[M];
#pragma omp for
  for(int i=0; i<N; i++){
    for(int j=0, j<M; j++)
      A[j] = (4*i+4*j) % M;
    for(int k=0; k<N; k++)
      X[i,k]=X[i,k-1] * A[A[(2*i+k)%M]];
  }
}
```



Reduction is Typically Easy To Recognize

If all the statements in which a scalar variable x appears are of the form $x \oplus = \text{exp}$, where x does not appear in exp and \oplus is associative then the cross-iteration RAWs on x can be resolved by:

- privatizing x initialized with the neutral element,
- computing the per-processor partial values of x ,
- reducing the x s across processors and with the initial value.

```
// compilation requires g++ -fopenmp ...
float x = 6.0;
#pragma omp parallel for reduction(+:x) private(i,j)
for(i=1; i<N; i++) {
    for(j=1; j<N; j++) {
        if ( A[i,j] >= 2.0 )    x += 2*A[i,j-1];
        else if( A[i,j] > 0.0 ) x += A[i-1,j+1];
    }
    if ( i % (j+1) == 3 ) x += A[i,i];
}
```



Scan and Segmented Scan Are Difficult!

Compilers cannot recognize and parallelize even simple scans:

- they raise a cross-iteration true dependency (RAW),
- they appear in a multitude of forms,
- hence they are difficult to analyze.

// What kind of scans are these?

1. `A[0] = B[0];`
 `for(i=1; i<N; i++) {`
 `A[i] = A[i-1] + B[i];`
 `}`
2. `acc = 0;`
 `for(i=0; i<N; i++){`
 `acc = acc xor i;`
 `A[i] = acc;`
 `}`
3. `for(j=0; j<M; j++)`
 `A[0,j] = B[0,j];`
 `for(i=1; i<N; i++) {`
 `for(j=0; j<M; j++)`
 `A[i,j] = A[i-1,j] + B[i,j];`
 `}`



Scan and Segmented Scan Are Difficult!

Compilers cannot recognize and parallelize even simple scans:

- they raise a cross-iteration true dependency (RAW),
- they appear in a multitude of forms,
- hence they are difficult to analyze.

// What kind of scans are these?

```
1. A[0] = B[0];
   for(i=1; i<N; i++) {
       A[i] = A[i-1] + B[i];
   }

2. acc = 0;
   for(i=0; i<N; i++){
       acc = acc xor i;
       A[i] = acc;
   }

3. for(j=0; j<M; j++)
    A[0,j] = B[0,j];
   for(i=1; i<N; i++) {
       for(j=0; j<M; j++)
           A[i,j] = A[i-1,j] + B[i,j];
   }
```

```
1. let A = scanInc (+) 0 B
```

```
2. let A = scanInc (xor) 0 [0..N-1]
```

```
3. let A = scanInc (\ a b -> zipWith (+) a b)
    (replicate M 0.0) B ≡
    let A = transpose $
        map (scanInc (+) 0.0) $
        transpose B
```



- 1 Direction-Vector Analysis
- 2 Block Tiling: Matrix Multiplication Case Study**
- 3 Coalesced Accesses: Matrix Transposition Case Study
- 4 Non-Trivial Synchronization: Histogram Case Study
- 5 Known Recurrences: Tridiagonal Solver Case Study
- 6 Imperative Context: Summarization of Array Indexes



Matrix Multiplication: Loop Strip Mining

```
DOALL i = 1, M, 1    // Parallel
  DOALL j = 1, N, 1  // Parallel
    float tmp = 0.0
    DO k = 1, U, 1 // Reduction
      tmp += A[i,k]*B[k,j]
    ENDDO
    C[i,j] = tmp;
  ENDDO
ENDDO
```

←Matrix Multiplication. Matrices:

- input A has M rows and U columns
- input B has U rows and N columns
- result C has M rows and N columns

Loops of indices i and j are parallel
(can be proved by direction vectors).

Accesses to A and B invariant to loops i and j \Rightarrow Block Tiling to optimize locality of reference!



Matrix Multiplication: Loop Strip Mining

```
DOALL i = 1, M, 1    // Parallel
  DOALL j = 1, N, 1  // Parallel
    float tmp = 0.0
    DO k = 1, U, 1 // Reduction
      tmp += A[i,k]*B[k,j]
    ENDDO
    C[i,j] = tmp;
  ENDDO
ENDDO
```

← Matrix Multiplication. Matrices:

- input A has M rows and U columns
- input B has U rows and N columns
- result C has M rows and N columns

Loops of indices i and j are parallel
(can be proved by direction vectors).

Accesses to A and B invariant to loops i and j \Rightarrow Block Tiling to optimize locality of reference!

First step: Strip Mining, always safe since the transformed loop executes the same instructions in the same order as the original loop:

```
DO i = 1, N, 1    // stride 1
  loop_body(i)
ENDDO
```

```
DO ii = 1, N, T    // stride T
  DO i = ii, MIN(ii+T-1,N), 1
    loop_body(i)
  ENDDO
ENDDO
```



Matrix Multiplication: Loop Interchange

After strip mining all loops with a tile of size T :

```
DOALL ii = 1, M, T
  DOALL i = ii, MIN(ii+T-1,M), 1      // loop
    DOALL jj = 1, N, T                // interchange
      DOALL j = jj, MIN(jj+T-1,N), 1
        float tmp = 0.0
        DO kk = 1, U, T
          DO k = kk, MIN(kk+T-1,U), 1
            tmp += A[i,k]*B[k,j]
          ENDDO ENDDO
        C[i,j] = tmp;
      ENDDO ENDDO ENDDO ENDDO
```

The second step is to apply loop interchange between the loops of indices i and jj . This is safe because loop i is parallel, hence it can always be interchanged inwards!



Matrix Multiplication: Summarizing Read Subscripts

After loop interchange we have a grid shape, as in CUDA:

```
DOALL ii = 1, M, T           // grid.y
  DOALL jj = 1, N, T         // grid.x
    DOALL i = ii, MIN(ii+T-1,M), 1 // block.y
      DOALL j = jj, MIN(jj+T-1,N), 1 // block.x
        float tmp = 0.0
        DO kk = 1, U, T
          DO k = kk, MIN(kk+T-1,U), 1
            tmp += A[i,k]*B[k,j]
          ENDDO
        ENDDO
      ENDDO
    ENDDO
  ENDDO
  C[i,j] = tmp;
ENDDO ENDDO ENDDO ENDDO
```

The third step is to summarize the subscripts of A and B read inside the loop of index k, for fixed ii, jj and kk (x:y denotes [x..y]):

- A subscripts [ii : MIN(ii+T-1,M), kk : MIN(kk+T-1,U)]
- B subscripts [kk : MIN(kk+T-1,U), jj : MIN(jj+T-1,N)]
- Summaries have size at most T^2 & independent on i, j, and k \Rightarrow CUDA-block threads cooperatively copy-in data to shared mem!



Block Tiled Matrix Multiplication CUDA Kernel

Shared memory padded with zeros to remove the branch from loop k!

```
DOALL ii = 1, M, T    // grid.y
  DOALL jj = 1, N, T // grid.x
    DOALL i = ii, MIN(ii+T-1,M), 1
      DOALL j = jj, MIN(jj+T-1,N), 1
        float tmp = 0.0
        DO kk = 1, U, T
          //we would like to copy
          //to shared memory here
          //& use it inside loop k
          DO k = kk, MIN(kk+T-1,U), 1
            tmp += A[i,k]*B[k,j]
          ENDDO
        ENDDO
      ENDDO
    ENDDO ENDDO ENDDO ENDDO
```



Block Tiled Matrix Multiplication CUDA Kernel

Shared memory padded with zeros to remove the branch from loop k!

```

DOALL ii = 1, M, T    // grid.y
  DOALL jj = 1, N, T  // grid.x
    DOALL i = ii, MIN(ii+T-1,M), 1
      DOALL j = jj, MIN(jj+T-1,N), 1
        float tmp = 0.0
        DO kk = 1, U, T
          //we would like to copy
          //to shared memory here
          //& use it inside loop k
          DO k = kk, MIN(kk+T-1,U), 1
            tmp += A[i,k]*B[k,j]
          ENDDO
        ENDDO
      ENDDO
    ENDDO ENDDO ENDDO ENDDO
  
```

```

__global__ void matMultTiledKer( ... ) {
  __shared__ T Ash[T][T], Bsh[T][T];
  int ii = blockIdx.y * T; //blockDim.x==T
  int jj = blockIdx.x * T; //blockDim.y==T
  int tidy = threadIdx.y, i = tidy+ii;
  int tidx = threadIdx.x, j = tidx+jj;
  float tmp = 0.0;

  for(int kk=0; kk<U; kk+=T) {
    Ash[tidy,tidx] = (i<M && kk+tidx<U) ?
                      A[i,kk+tidx] : 0.0 ;
    Bsh[tidy,tidx] = (j<N && kk_tidx<U) ?
                      B[kk+tidy,j] : 0.0 ;

    __syncthreads();
    for(int k=0; k<T; k++) {
      tmp += Ash[tidy][k] * Bsh[k][tidx]
    } __syncthreads();
  } if (i<M && j<N) C[i,j] = tmp;
}
  
```

A global memory access amortized by (T-1) shared memory accesses.



Measuring GFlops Performance

Sequential matrix multiplication $\sim 2 \times M \times N \times U$ floating point operations. What is the GFlops performance of our implementation?

```
// CPU code
int dimy = ceil( ((float)M) / T );
int dimx = ceil( ((float)N) / T );
dim3 block(T,T,1), grid(dimx,dimy,1);

unsigned long int elapsed;
struct timeval t_start,t_end,t_diff;
gettimeofday(&t_start, NULL);

// ignoring generic shared mem problems
matMultTiledKer<T><<<grid, block>>> (d_A, d_B, d_C, U, M, N);

gettimeofday(&t_end, NULL);
timeval_subtract(&t_diff,
                 &t_end,&t_start);
elapsed=(t_diff.tv_sec*1e6 +
         t_diff.tv_usec);
double flops = 2.0 * M * N * U;
double gigaFlops=(flops*1.0e-3f) /
                 elapsed;

template <int T> // KERNEL
__global__ void matMultTiledKer( ... ) {
    __shared__ float Ash[T][T], Bsh[T][T];
    int ii = blockIdx.y * T; //blockDim.x==T
    int jj = blockIdx.x * T; //blockDim.y==T
    int tidy = threadIdx.y, i = tidy+ii;
    int tidx = threadIdx.x, j = tidx+jj;
    float tmp = 0.0;

    for(int kk=0; kk<U; kk+=T) {
        Ash[tidy,tidx] = (i<M && kk+tidx<U) ?
                          A[i,kk+tidx] : 0.0 ;
        Bsh[tidy,tidx] = (j<N && kk_tidy<U) ?
                          B[kk+tidy,j] : 0.0 ;

        __syncthreads();
        for(int k=0; k<T; k++) {
            tmp += Ash[tidy][k] * Bsh[k][tidx];
        } __syncthreads();
    } if (i<M && j<N) C[i,j] = tmp;
}
```



- 1 Direction-Vector Analysis
- 2 Block Tiling: Matrix Multiplication Case Study
- 3 Coalesced Accesses: Matrix Transposition Case Study
- 4 Non-Trivial Synchronization: Histogram Case Study
- 5 Known Recurrences: Tridiagonal Solver Case Study
- 6 Imperative Context: Summarization of Array Indexes



Matrix Transposition: Motivation

```
// Non-Coalesced Memory Access
// Transposition to coalesce it ⇒
DOALL i = 0 to N-1 // parallel
  tmpB = A[i,0] * A[i,0]
  B[i,0] = tmpB
DO j = 1, 63 // sequential
  tmpA = A[i, j]
  accum = tmpB*tmpB + tmpA*tmpA
  B[i,j] = accum
  tmpB = accum
ENDDO
ENDDO

A' = transpose(A)
DOALL i = 0 to N-1 // parallel
  tmpB = A'[0,i] * A'[0,i]
  B'[0,i] = tmpB
DO j = 1, 63 // sequential
  tmpA = A'[j, i]
  accum = tmpB*tmpB + tmpA*tmpA
  B'[j, i] = accum
  tmpB = accum
ENDDO
ENDDO
B = transpose(B')
```

The transformed program performs about twice the number of accesses to global memory than the original.

But exhibits only coalesced accesses!

What else could we have done to achieve the same effect?



Matrix Transposition: Motivation

```
// Non-Coalesced Memory Access
// Transposition to coalesce it ⇒
DOALL i = 0 to N-1 // parallel
  tmpB = A[i,0] * A[i,0]
  B[i,0] = tmpB
DO j = 1, 63 // sequential
  tmpA = A[i, j]
  accum = tmpB*tmpB + tmpA*tmpA
  B[i,j] = accum
  tmpB = accum
ENDDO
ENDDO

A' = transpose(A)
DOALL i = 0 to N-1 // parallel
  tmpB = A'[0,i] * A'[0,i]
  B'[0,i] = tmpB
DO j = 1, 63 // sequential
  tmpA = A'[j, i]
  accum = tmpB*tmpB + tmpA*tmpA
  B'[j, i] = accum
  tmpB = accum
ENDDO
ENDDO
B = transpose(B')
```

The transformed program performs about twice the number of accesses to global memory than the original.

But exhibits only coalesced accesses!

What else could we have done to achieve the same effect?

Loop Interchange. Does it work in general?



Transposition: Strip Mining, Interchange & Kernel

```
//Both loops are parallel
//Strip mining & interchange⇒
for(i = 0; i < rowsA; i++) {
    for(j = 0; j < colsA; j++) {
        trA[j*rowsA+i] = A[i*colsA+j];
    } }
```

```
for(ii=0; ii<rowsA; ii+=T) {
    for(jj=0; jj<colsA; jj+=T) {
        for(i=ii; i<min(ii+T,rowsA); i++) {
            for(j=jj; j<min(jj+T,colsA); j++) {
                trA[j*rowsA+i] = A[i*colsA+j];
            } } } }
```

```
__global__ void matTranspose(
    float* A, float* trA,
    int rowsA, int colsA ) {
    __shared__ float tile[T][T];
    int tidx = threadIdx.x;
    int tidy = threadIdx.y;
    int j = blockIdx.x*T + tidx;
    int i = blockIdx.y*T + tidy;
    if( j < colsA && i < rowsA )
        tile[tidy][tidx] = A[i*colsA+j];
    __syncthreads();
    i = blockIdx.y*T + threadIdx.x;
    j = blockIdx.x*T + threadIdx.y;
    if( j < colsA && i < rowsA )
        trA[j*rowsA+i] = tile[tidx][tidy];
}
```

- Trick is to write the element of the symmetric thread in the same block.
- What is the problem?



Transposition: Strip Mining, Interchange & Kernel

```
//Both loops are parallel
//Strip mining & interchange⇒
for(i = 0; i < rowsA; i++) {
    for(j = 0; j < colsA; j++) {
        trA[j*rowsA+i] = A[i*colsA+j];
    } }
```

```
for(ii=0; ii<rowsA; ii+=T) {
    for(jj=0; jj<colsA; jj+=T) {
        for(i=ii; i<min(ii+T,rowsA); i++) {
            for(j=jj; j<min(jj+T,colsA); j++) {
                trA[j*rowsA+i] = A[i*colsA+j];
            } } } }
```

```
__global__ void matTranspose(
    float* A, float* trA,
    int rowsA, int colsA ) {
    __shared__ float tile[T][T];
    int tidx = threadIdx.x;
    int tidy = threadIdx.y;
    int j = blockIdx.x*T + tidx;
    int i = blockIdx.y*T + tidy;
    if( j < colsA && i < rowsA )
        tile[tidy][tidx] = A[i*colsA+j];
    __syncthreads();
    i = blockIdx.y*T + threadIdx.x;
    j = blockIdx.x*T + threadIdx.y;
    if( j < colsA && i < rowsA )
        trA[j*rowsA+i] = tile[tidx][tidy];
}
```

- Trick is to write the element of the symmetric thread in the same block.
- What is the problem?
- Number of shared memory banks typically 16 or 32.
- T is also either 16 or 32 ⇒
- 16 consecutive threads will read the same memory bank at the same time.
- Solution: `tile[T][T+1];`



Lessons Learned So Far

- Tiled transposition is less than $3\times$ faster than the naive version,
- but the motivating example runs more than $8\times$ faster when transposition coalesces accesses to arrays A and B. **Why?**



Lessons Learned So Far

- Tiled transposition is less than $3\times$ faster than the naive version,
- but the motivating example runs more than $8\times$ faster when transposition coalesces accesses to arrays A and B. **Why?**
- **Better to eliminate rather than hide latency.** Impact of hardware multi-threading limited by the amount of available resources!
- Generic Array of Tuples using shared memory:



Lessons Learned So Far

- Tiled transposition is less than $3\times$ faster than the naive version,
- but the motivating example runs more than $8\times$ faster when transposition coalesces accesses to arrays A and B. **Why?**
- **Better to eliminate rather than hide latency.** Impact of hardware multi-threading limited by the amount of available resources!
- Generic Array of Tuples using shared memory: requires a class with a “volatile” assignment operator [Esben Skaarup].
- I believe it is safe to ignore the warnings; they can be eliminated at the expense of writing awkward code.
- Shared-memory of generic type [Aske Dorge]:



Lessons Learned So Far

- Tiled transposition is less than $3\times$ faster than the naive version,
- but the motivating example runs more than $8\times$ faster when transposition coalesces accesses to arrays A and B. **Why?**
- **Better to eliminate rather than hide latency.** Impact of hardware multi-threading limited by the amount of available resources!
- Generic Array of Tuples using shared memory: requires a class with a “volatile” assignment operator [Esben Skaarup].
- I believe it is safe to ignore the warnings; they can be eliminated at the expense of writing awkward code.
- Shared-memory of generic type [Aske Dorge]: empirically, ~ 8 words of shared memory per thread does not degrade perform.
- Simple solution: declare in kernel a shared memory array of type `char`, then cast it to the generic type. Set the size when calling the kernel to `block_size*32` bytes.



Constant (Read-Only) Memory in CUDA

- 64KB of `__constant__` memory on device (global/slow), cached in each multiprocessor, e.g., 8KB (fast).
- May reduce the required memory bandwidth:
 - if found in cache, then no extra traffic,
 - if a (half) warp accesses the same location and misses in cache \Rightarrow only one request is sent and the result is broadcast back to all,
 - **serialized** accesses if a warp of threads read different locations!
 - latency can range from one to hundreds of cycles.
- **Best Use:** when an entire block accesses the same location in the same SIMD instruction: even on a miss, the first warp brings the data in cache @ minimal traffic, the rest find it in cache.

```
// C in __constant__ memory: Good!
DO i = 1, N, 1 // grid
  DO j = 1, M, 1 // block(s)
    A[i,j] = A[i,j] % C[i]
  ENDDO

// C in __constant__ memory: Bad!
DO i = 1, N, 1 // grid
  DO j = 1, M, 1 // block(s)
    A[i,j] = A[i,j] % C[j]
  ENDDO
// Either global memory or loop interchange
```



- 1 Direction-Vector Analysis
- 2 Block Tiling: Matrix Multiplication Case Study
- 3 Coalesced Accesses: Matrix Transposition Case Study
- 4 Non-Trivial Synchronization: Histogram Case Study**
- 5 Known Recurrences: Tridiagonal Solver Case Study
- 6 Imperative Context: Summarization of Array Indexes



Histogram: Problem Statement

Sequential Pseudocode

&

Notation/Parameters:

```
for(int i = 0; i < N; i++ ) {  
    (ind, val) = fun(image,i);  
    H[ ind ] += val;  
}
```

H = histogram
B = sizeof(H), e.g., $4 \times 4 \dots 64 \times 64$
N = sizeof(image), e.g., $256 \times 256 \times 256$
P = degree of parallelism (# of cores)

Typically, the additions to a histogram are 1, but in this case two locations of the histogram are incremented with `val` and `1.0-val` (not shown).

`ind` values are unstructured and may generate data races if loop iterations are executed out of order, i.e., in parallel.

However the code looks like a reduction on arrays \Rightarrow can be parallelized.



Histogram: OpenMP Code

One Histogram/Thread

&

Reduce Across Histograms

```
int  P  = omp_get_num_threads();
REAL* pH = new REAL[P*B];
#pragma omp parallel
{ int  th_id = omp_get_thread_num();
  REAL* privH = pH + th_id * B;
  for(int i=0; i<B; i++) privH[i]=0.0;

#pragma omp for schedule(static)
  for( int i=0; i < N; i++ ) {
    (ind, val) = fun(image,i);
    privH[ ind ] += val;
  } // end par for
```

```
// Assuming B big enough
pH' = parallel_transpose(pH);

#pragma omp for schedule(static)
  for( int i=0; i<B; i++ ) {
    for( int p=0; p<P; p++ ) {
      H[i] += pH'[i*P + p];
    }
  } // end par for
} // end OMP parallel region
```

Work Efficient: $O(P \times B) \leq O(N)$. Typically only the outer loop of the (image) nest is executed in parallel, since P is relatively small.

The alternative is to not privatize & reduce, but instead to provide lock-free (fine-grain) synchronization for each index in the histogram.



What Is Different for GPGPU Parallelization ?

Differences:

- Quite a few more cores, hence need to exploit all levels of parallelism,
- No (illusion of) constant-time random-memory access, e.g., global memory accesses are up to two order of magnitude more expensive than (scarce) local memory,
- Synchronization primitives limited to local workgroup, i.e., barriers, and can simulate CAS-like instructions at warp level.



Intuition

- If possible, keep histograms in local memory, e.g., about 16 “local” floats per thread:
- If $B \leq 16$ then 1 histogram is maintained per thread,
- if $B = 32$ then two threads build one histogram, using compare-and-swap CAS like locking.
- In general if $16 < B \leq 32 * 8 = 512$ then up to 32 (warp size) threads cooperate at building a histogram.
- **Difficulty:** up to $B = 16 \times 128$ one workgroup builds a histogram; unfortunately no efficient CAS-like sync. at workgroup level.
- **Difficulty:** for $B \geq 64 * 64$ need to keep the histogram in global memory: prohibitively expensive due to irregular accesses.

For simplicity use fix $L = 128$ the size of the local workgroup,
 $G = 512$ the number of workgroups, and unroll the loop: $N/(L \times G)$.
Work Efficient $\sim B \leq N / G$.



When L / WARP Histograms Fit in Shared Memory

CUDA Pseudocode; , C is the number of cooperating threads

```

__shared__ REAL sh_hist[ 16 * L ]; // L is the size of the block.
int tid = threadIdx.x;
// init local histogram(s)
for( t = tid; t < 16*L; t += L )
    sh_hist[ t ] = 0.0;
__syncthreads();

// 2) each thread executes U iterations sequentially.
u = blockIdx.x * U * L + tid;
for( t = u; t < u+(U*L); t += L ) {
    (ind, val) = fun(image,i);
    // we work with the transposed sh_hist
    ind = ( ind * L / C ) + ( tid / C );
    if ( C > 1 )
        raw_cas_update( sh_hist, ind, val );
    else sh_hist[ ind ] += val;
} __syncthreads();

// 3) reduce locally across histograms
segm_scan_reg_block( sh_hist, B * L / C, L / C );
// 4) commit the resulted local histogram to global storage ...

```



CAS-like Synchronization

Function raw_cas_update

```
void raw_cas_update(volatile REAL* sh_hist, ulong ind, REAL val) {  
    REAL acc, id = (REAL)threadIdx.x;  
    if (val == 0.0) return;  
    bool repeat = true;  
    while(repeat) {  
        acc = sh_hist[ind];  
        sh_hist[ind] = id;  
        if( sh_hist[ ind ] == id ) {  
            sh_hist[ ind ] = acc + val;  
            repeat = false;  
        }  
    }  
}
```

Bug in NVIDIA implem of OPENCL: if barriers are placed inside while loop the execution does not deadlocks but incorrect result!

How Do Results Look Like?



- ① Direction-Vector Analysis
- ② Block Tiling: Matrix Multiplication Case Study
- ③ Coalesced Accesses: Matrix Transposition Case Study
- ④ Non-Trivial Synchronization: Histogram Case Study
- ⑤ Known Recurrences: Tridiagonal Solver Case Study
- ⑥ Imperative Context: Summarization of Array Indexes



TRIDAG Motivation

Stochastic Volatility Calibration: given a set of observed prices of contracts, identify a model of such prices, as a function of volatility (unknown), time and strikes (known), and some other (unobserved) parameters.

Volatility is modeled as a system of continuous differential equations and solved via Crank-Nicolson finite-difference method:

$$\frac{\partial f}{\partial t}(x, t) + \mu(x, t) \frac{\partial f}{\partial x}(x, t) + \frac{1}{2} \sigma(x, t)^2 \frac{\partial^2 f}{\partial x^2}(x, t) - r(x, t) f(x, t) = 0 \quad (1)$$

with terminal condition $f(x, T) = F(x)$, $x \in \mathcal{S}$, where F is known.

It comes down to solving this equation for many instances of μ, σ, r .



Explicit Finite Difference Approach

- $\Delta x = (x_J - x_1)/J$, $\Delta t = (t_N - t_1)/N$
- Uses Parametrizations:
 - $D_x f_{j,n} = \frac{f_{j+1,n} - f_{j-1,n}}{2\Delta x}$, $D_x^2 f_{j,n} = \frac{f_{j+1,n} - 2f_{j,n} + f_{j-1,n}}{(\Delta x)^2}$
 - $D_t^- f_{j,n} = \frac{f_{j,n} - f_{j,n-1}}{\Delta t}$
- Using this discretization in the differential equation yields:
$$f_{j,n-1} = \alpha_{j,n} f_{j-1,n} + \beta_{j,n} f_{j,n} + \gamma_{j,n} f_{j+1,n}$$
- where $f_{j,N}$ are known $\forall j \in \{1 \dots J\}$ and we aim to find $f_{j,1}, \forall j$.
- Trivial Parallel Algorithm Depth $O(T)$, Work $O(JT)$
- However, mathematical reasoning requires $N \gg J$, i.e., a very fine-grained time discretization \Rightarrow deep depth!



Implicit Finite Difference Approach

- $\Delta x = (x_J - x_1)/J$, $\Delta t = (t_N - t_1)/N$
- Uses Parametrizations:
 - $D_x f_{j,n} = \frac{f_{j+1,n} - f_{j-1,n}}{2\Delta x}$, $D_x^2 f_{j,n} = \frac{f_{j+1,n} - 2f_{j,n} + f_{j-1,n}}{(\Delta x)^2}$
 - $D_t^+ f_{j,n} = \frac{f_{j,n+1} - f_{j,n}}{\Delta t}$
- Using this discretization in the differential equation yields:

$$f_{j,n+1} = a_{j,n}f_{j-1,n} + b_{j,n}f_{j,n} + c_{j,n}f_{j+1,n}$$
- where $f_{j,N}$ are known $\forall j \in \{1 \dots J\}$ and we aim to find $f_{j,1}, \forall j$,
- meaning that we know $f_{j,n+1} \forall j$ and want to compute $f_{j,n} \forall j$
- Requires solving a tridiagonal system (TRIDAG) at every time step; non-trivial to parallelize!
- However, mathematical reasoning requires $N \sim J$, i.e., the depth is reduced and degree of parallelism increased.



TRIDAG Problem Statement

Needed: Parallel Algorithm for Computing X :

Given A and D , find X such that $A * X = D$,
 where $A \in \mathbb{M}^{n \times n}$ is **tridiagonal**, and X and D vectors of size n .

A										$*$	X	$=$	D						
	b_0	c_0	0	0	0				x_0				d_0					
	a_0	b_1	c_1	0	0				x_1				d_1					
	0	a_1	b_2	c_2	0	0				x_2				d_2				
											$*$			$=$		
	0	0	...	0	a_{n-3}	b_{n-2}	c_{n-2}				x_{n-2}				d_{n-2}				
	0	0	...	0	0	a_{n-2}	b_{n-1}				x_{n-1}				d_{n-1}				

At every step in the time series/discretization (sequential), we know matrix A and vector D and need to compute in parallel vector X .



High-Level Solution (Sequential and Parallel)

Problem is to Find $X = [x_0, \dots, x_{n-1}]^T$ such that:

$$\begin{array}{cccccccc|cccc|}
 & A & & * & X & = & D & & & & & \\
 | & b_0 & c_0 & 0 & 0 & & \dots & 0 & | & | & x_0 & | & | & d_0 & | \\
 | & a_0 & b_1 & c_1 & 0 & & \dots & 0 & | & | & x_1 & | & | & d_1 & | \\
 | & 0 & a_1 & b_2 & c_2 & 0 & & \dots & 0 & | & | & x_2 & | & | & d_2 & | \\
 | & \dots & \dots & \dots & \dots & \dots & \dots & \dots & | & * & | & \dots & | & = & | & \dots & | \\
 | & 0 & 0 & \dots & 0 & a_{n-3} & b_{n-2} & c_{n-2} & | & | & x_{n-2} & | & | & d_{n-2} & | \\
 | & 0 & 0 & \dots & 0 & 0 & a_{n-2} & b_{n-1} & | & | & x_{n-1} & | & | & d_{n-1} & |
 \end{array}$$

- STEP 1: Compute the LU decomposition of A , i.e., $A = L * U$, where L and U are lower and upper-diagonal matrices $\in \mathbb{M}^{n \times n}$.
- STEP 2: Solve $L * Y = D$, i.e., compute partial solution Y .
- STEP 3: Solve $U * X = Y$, i.e., compute problem solution X .
- Each of the three steps needs to be parallel (of depth $O(\log n)$)!



High-Level Solution (Sequential and Parallel)

Problem is to Find $X = [x_0, \dots, x_{n-1}]^T$ such that:

$$\begin{array}{cccccccc|cccc|}
 & A & & & & & & & * & X & = & D \\
 | & b_0 & c_0 & 0 & 0 & & & 0 & | & x_0 & | & d_0 & | \\
 | & a_0 & b_1 & c_1 & 0 & & & 0 & | & x_1 & | & d_1 & | \\
 | & 0 & a_1 & b_2 & c_2 & 0 & & 0 & | & x_2 & | & d_2 & | \\
 | & \dots & \dots & \dots & \dots & \dots & \dots & \dots & | & * & | & \dots & | \\
 | & 0 & 0 & \dots & 0 & a_{n-3} & b_{n-2} & c_{n-2} & | & x_{n-2} & | & d_{n-2} & | \\
 | & 0 & 0 & \dots & 0 & 0 & a_{n-2} & b_{n-1} & | & x_{n-1} & | & d_{n-1} & |
 \end{array}$$

```

void tridag( REAL* a, REAL* b, REAL* c, REAL* d, int n, REAL* y, REAL* u ){
    double beta;
    y[0] = d[0];    u[0] = b[0];
    for(int i=1; i<n; i++) {        // distribute with
        beta = a[i] / u[i-1];      // direction vcts!
        u[i] = b[i] - beta*c[i-1];
        y[i] = d[i] - beta*y[i-1];
    }
    y[n-1] = y[n-1]/u[n-1];
    for(int i=n-2; i>=0; i--) {
        y[i] = (y[i] - c[i]*y[i+1]) / u[i];
    } }

```



LU-Decomposition Recurrences

Writing the LU decomposition: $A = L * U$, where $L, U \in \mathbb{M}^{n \times n}$

$$\begin{array}{c|c|c|c|c|c}
 A & = & L & * & U & \\
 \hline
 \begin{array}{l}
 | \text{ } b_0 \text{ } c_0 \text{ } 0 \text{ } 0 \text{ } \dots \text{ } 0 \text{ } | \\
 | \text{ } a_0 \text{ } b_1 \text{ } c_1 \text{ } 0 \text{ } \dots \text{ } 0 \text{ } | \\
 | \text{ } \dots \text{ } \dots \text{ } \dots \text{ } c_{n-2} \text{ } | \\
 | \text{ } 0 \text{ } \text{ } 0 \text{ } \dots \text{ } 0 \text{ } a_{n-2} \text{ } b_{n-1} \text{ } |
 \end{array}
 & = &
 \begin{array}{l}
 | \text{ } 1 \text{ } 0 \text{ } 0 \text{ } \dots \text{ } 0 \text{ } | \\
 | \text{ } m_0 \text{ } 1 \text{ } 0 \text{ } \dots \text{ } 0 \text{ } | \\
 | \text{ } \dots \text{ } \dots \text{ } \dots \text{ } \text{ } | \\
 | \text{ } 0 \text{ } \dots \text{ } 0 \text{ } m_{n-2} \text{ } 1 \text{ } |
 \end{array}
 & * &
 \begin{array}{l}
 | \text{ } u_0 \text{ } c_0 \text{ } 0 \text{ } \dots \text{ } 0 \text{ } | \\
 | \text{ } 0 \text{ } u_1 \text{ } c_1 \text{ } \dots \text{ } 0 \text{ } | \\
 | \text{ } \dots \text{ } \dots \text{ } u_{n-2} \text{ } c_{n-2} \text{ } | \\
 | \text{ } 0 \text{ } \dots \text{ } 0 \text{ } u_{n-1} \text{ } |
 \end{array}
 \end{array}$$

LU-decomposition Recurrences:

I. First, compute u_i , $i \in \{0 \dots n-1\}$:

I.a) $u_0 = b_0$

I.b) $u_i = b_i - (a_{i-1} * c_{i-1}) / u_{i-1}$, $i \in \{1 \dots n-1\}$

II. Knowing the u 's, we can easily compute the m 's:

$$m_i = a_i / u_i, i \in \{0 \dots n-2\}$$

The Difficult Recurrence of LU-Decomposition I.b)

Parallel Algorithm: **scan** with $\mathbb{M}^{2 \times 2}$ multiplication operator

Recurrence:

$$\text{I.a)} \quad u_0 = b_0$$

$$\text{I.b)} \quad u_i = b_i - (a_{i-1} * c_{i-1}) / u_{i-1}, \quad i \in \{1 \dots n-1\}$$

1) Substitute in I.b) $u_i \leftarrow q_{i+1} / q_i$,
 where $q_1 = b_0$ and $q_0 = 1.0$, i.e., $u_0 == b_0$ still holds! We obtain:

$$\text{I.c)} \quad \begin{vmatrix} q_{i+1} \\ q_i \end{vmatrix} = \begin{vmatrix} b_i & -a_{i-1} * c_{i-1} \\ 1.0 & 0.0 \end{vmatrix} * \begin{vmatrix} q_i \\ q_{i-1} \end{vmatrix} \quad i \in \{1 \dots n-1\}$$

2) Denoting the above 2×2 matrix by $S_i \in \mathbb{M}^{2 \times 2}$ we obtain by induction:

$$\text{I.d)} \quad \begin{vmatrix} q_{i+1} & q_i \end{vmatrix}^T = (S_i * S_{i-1} * S_1) * \begin{vmatrix} b_0 & 1.0 \end{vmatrix}^T$$

3) We can compute all such matrices, i.e., $S_i * S_{i-1} * S_1$ with $i \in \{1..n-1\}$,
 by applying **scanInc** with the (associative) matrix-multiplication operator (*)

$$\text{I.e)} \quad \text{matrices} = \text{scanInc} (*) [S_{n-1}, S_{n-2}, \dots, S_1]$$

4) Finally, we compute each $[q_{i+1}, q_i]^T$ by multiplying its corresponding
 matrix with $[b_0, 1.0]^T$, and compute $u_i = q_{i+1} / q_i$. Both are **PARALLEL**:

$$\text{I.f)} \quad \text{q_pairs} = \text{map} (\text{matVectMult} (b_0, 1.0)) \text{ matrices}$$

$$\text{I.g)} \quad u = \text{map} (\backslash (x,y) \rightarrow x/y) \text{ q_pairs}$$



STEP 2 & 3: Solving For/Backward Recurrences

STEP 2: Parallel Algorithm for $L * Y = D$ (forward recurrence)

- 1) By doing the arithmetics, it comes down to solving the recurrence:
 - II.a) $y_0 = d_0$
 - II.b) $y_i = d_i - m_{i-1} * y_{i-1}, \quad i \in 1 .. n-1$
- 2) We shall use **scanInc** with linear-function-composition operator:
 - i) Consider $f_1(x) = a_1 + b_1*x$ and $f_2(x) = a_2 + b_2*x$
 - ii) $(f_1 \diamond f_2)(x) = (a_2+b_2*a_1) + (b_1*b_2)*x$
 - iii) We represent such a function f as a pair (a, b) , and implement function application via operator: $\text{apply } x \ (a, b) = a + b*x$
- 3) Representing $f_i = (d_i, -m_{i-1}), \quad i \in 1 .. n-1$
 - i) It follows by induction that $y_i = (f_i \diamond f_{i-1} \diamond \dots \diamond f_1)(d_0)$
 - ii) We do a **scanInc** with \diamond and neutral element $(0.0, 1.0)$:

$$f_out = \text{scanInc } (\diamond) \ (0.0, 1.0) \ [f_1, \dots, f_{n-1}]$$
 - iii) Finally, apply the f_out 's to d_0 to compute all $y_i, \quad i \in \{1 .. n-1\}$

$$y = \text{map } (\text{apply } d_0) \ f_out$$

STEP 3: Solving $U * X = Y$ is similar to STEP 2 (but *backwards*):

III.a) $x_{n-1} = y_{n-1}/u_{n-1}$ b) $x_i = y_i/u_i - c_i*x_{i+1}/u_i, \quad i \in \{n-2..0\}$



- 1 Direction-Vector Analysis
- 2 Block Tiling: Matrix Multiplication Case Study
- 3 Coalesced Accesses: Matrix Transposition Case Study
- 4 Non-Trivial Synchronization: Histogram Case Study
- 5 Known Recurrences: Tridiagonal Solver Case Study
- 6 Imperative Context: Summarization of Array Indexes




Interprocedural Summarization of Array Indexes

Independence-Summary Simple Example

```

DO i = 1, N
  A(i+100) = ...
  IF (x > 0) THEN
    ... = A(i)
  ENDIF
ENDDO

```



- Techniques that analyze read-write pairs of accesses become very conservative on larger loops with non-trivial control flow.
- *Alternative:* inter-procedural summarization + model loop independence via an equation on summaries of shape $S = \emptyset$
- Decrease overhead by extracting lightweight predicates that prove independence at runtime, e.g., $x \leq 0 \vee N < 100$.

Show Calculix from SPEC2006!



Building RO, RW, WF Summaries Interprocedurally

Summaries (RO, RW, WF) are

- constructed via a bottom-up parse of the CALL and CD graphs,
- structural data-flow equations dictate how to compose consecutive regions, aggregate/translate across loops/callsites, ...



Building RO, RW, WF Summaries Interprocedurally

Summaries (RO, RW, WF) are

- constructed via a bottom-up parse of the CALL and CD graphs,
- structural data-flow equations dictate how to compose consecutive regions, aggregate/translate across loops/callsites, ...

Simplified solvh_do20 from dyfesm

```

DO i = 1, N
  CALL geteu (XE(IA(i)), NP, SYM)
  CALL matmul(XE(IA(i)), NS)
ENDDO

SUBROUTINE matmul(XE, NS)
  INTEGER NS, XE(*)
  DO j = 1, NS
    ... = XE(j) ...
    XE(j) = ...
  ENDDO
END

SUBROUTINE geteu(XE, NP, SYM)
  INTEGER NP, SYM, XE(16, *)
  IF (SYM .NE. 1) THEN
    DO i = 1, NP
      DO j = 1, 16
        XE(j, i) = ...
      ENDDO
    ENDDO
  ENENDIF
END

```



Summarizing Subroutine geteu

WF summary for geteu; $RO_{geteu} = RW_{geteu} = \emptyset$

```

 $S_{geteu}$   SUBROUTINE geteu(XE, NP, SYM)
           INTEGER NP, SYM, XE(16, *)
 $S_{If}$       IF (SYM .NE. 1) THEN
 $S_{Li}$         DO i = 1, NP
 $S_{Lj}$           DO j = 1, 16
 $S_{WF}$             XE(j, i) = ...
                ENDDO
            ENDDO
        ENDIF
    END
  
```

$$WF_{S_{WF}}^{XE} = \{16 * i + j - 1\}$$

- Loop Aggregation uses (intuitively) interval arithmetic:
- Loop i : $\{16 * i + j - 1 \mid j \in \{1..16\}\} \rightarrow 16 * i + [0, 15]$
- Loop j : $\{16 * i + [0, 15] \mid i \in \{1..NP\}\} \rightarrow [0, 16 * NP - 1]$
- Branches introduce predicated nodes, e.g., $WF_{S_{if}}^{XE} = WF_{S_{geteu}}^{XE}$



Summarizing Subroutine geteu

WF summary for geteu; $RO_{geteu} = RW_{geteu} = \emptyset$

```

 $S_{geteu}$   SUBROUTINE geteu(XE, NP, SYM)
           INTEGER NP, SYM, XE(16, *)
 $S_{If}$       IF (SYM .NE. 1) THEN
 $S_{Li}$         DO i = 1, NP
 $S_{Lj}$           DO j = 1, 16
 $S_{WF}$             XE(j, i) = ...
                ENDDO
            ENDDO
        ENDIF
    END

```

$WF_{S_{Lj}}^{XE} = 16 * i + [0, 15]$
 $WF_{S_{WF}}^{XE} = \{16 * i + j - 1\}$

- Loop Aggregation uses (intuitively) interval arithmetic:
- Loop i : $\{16 * i + j - 1 \mid j \in \{1..16\}\} \rightarrow 16 * i + [0, 15]$
- Loop j : $\{16 * i + [0, 15] \mid i \in \{1..NP\}\} \rightarrow [0, 16 * NP - 1]$
- Branches introduce predicated nodes, e.g., $WF_{S_{if}}^{XE} = WF_{S_{geteu}}^{XE}$



Summarizing Subroutine geteu

WF summary for geteu; $RO_{geteu} = RW_{geteu} = \emptyset$

```

 $S_{geteu}$   SUBROUTINE geteu(XE, NP, SYM)
           INTEGER NP, SYM, XE(16, *)
 $S_{If}$       IF (SYM .NE. 1) THEN
 $S_{Li}$         DO i = 1, NP
 $S_{Lj}$           DO j = 1, 16
 $S_{WF}$             XE(j, i) = ...
                ENDDO
            ENDDO
        ENDIF
    END

```

$$WF_{S_{Li}}^{XE} = [0, 16 * NP - 1]$$

$$WF_{S_{Lj}}^{XE} = 16 * i + [0, 15]$$

$$WF_{S_{WF}}^{XE} = \{16 * i + j - 1\}$$

- Loop Aggregation uses (intuitively) interval arithmetic:
- Loop i : $\{16 * i + j - 1 \mid j \in \{1..16\}\} \rightarrow 16 * i + [0, 15]$
- Loop j : $\{16 * i + [0, 15] \mid i \in \{1..NP\}\} \rightarrow [0, 16 * NP - 1]$
- Branches introduce predicated nodes, e.g., $WF_{S_{if}}^{XE} = WF_{S_{geteu}}^{XE}$



Summarizing Subroutine geteu

WF summary for geteu; $RO_{geteu} = RW_{geteu} = \emptyset$

S_{geteu} SUBROUTINE geteu(XE, NP, SYM) INTEGER NP, SYM, XE(16, *) S_{IF} IF (SYM .NE. 1) THEN S_{Li} DO i = 1, NP S_{Lj} DO j = 1, 16 S_{WF} XE(j, i) = ... ENDDO ENDDO ENDIF END	$WF_{S_{IF}}^{XE} = \begin{matrix} (SYM \neq 1) \\ \downarrow \\ [0, 16 * NP - 1] \end{matrix}$ $WF_{S_{Li}}^{XE} = [0, 16 * NP - 1]$ $WF_{S_{Lj}}^{XE} = 16 * i + [0, 15]$ $WF_{S_{WF}}^{XE} = \{16 * i + j - 1\}$
---	--

- Loop Aggregation uses (intuitively) interval arithmetic:
- Loop i : $\{16 * i + j - 1 \mid j \in \{1..16\}\} \rightarrow 16 * i + [0, 15]$
- Loop j : $\{16 * i + [0, 15] \mid i \in \{1..NP\}\} \rightarrow [0, 16 * NP - 1]$
- Branches introduce predicated nodes, e.g., $WF_{S_{if}}^{XE} = WF_{S_{geteu}}^{XE}$



Summarizing Subroutine matmul

RW summary for matmul; $RO_{matmul} = WF_{matmul} = \emptyset$

S_{matmul} SUBROUTINE matmul(XE, NS)

INTEGER NS, XE(*)

S_{loop} DO j = 1, NS

S_{RO} ... = XE(j) ...

S_{WF} XE(j) = ...

ENDDO

END

$$RO_{S_{RO}}^{XE} = \{j-1\}$$

$$WF_{S_{WF}}^{XE} = \{j-1\}$$

- Composing read-only RO_{S_1} and write-first WF_{S_2} regions:
- $RO = RO_{S_1} - WF_{S_2}$, $WF = WF_{S_2} - RO_{S_1}$, $RW = RO_{S_1} \cap WF_{S_2}$
- In our case $RO = \emptyset$, $WF = \emptyset$, $RW = \{j-1\}$
- Over loop DO j: $RO_{loop} = \emptyset$, $WF_{loop} = \emptyset$, $RW_{loop} = [0, NS-1]$



Summarizing Subroutine matmul

RW summary for matmul; $RO_{matmul} = WF_{matmul} = \emptyset$

S_{matmul} SUBROUTINE matmul(XE, NS)

INTEGER NS, XE(*)

S_{loop} DO j = 1, NS

S_{RO} ... = XE(j) ...

S_{WF} XE(j) = ...

ENDDO

END

$$S_{RO} \diamond S_{WF} = \{\emptyset, \emptyset, RW = \{j-1\}\}$$

$$RO_{S_{RO}}^{XE} = \{j-1\}$$

$$WF_{S_{WF}}^{XE} = \{j-1\}$$

- Composing read-only RO_{S_1} and write-first WF_{S_2} regions:
- $RO = RO_{S_1} - WF_{S_2}$, $WF = WF_{S_2} - RO_{S_1}$, $RW = RO_{S_1} \cap WF_{S_2}$
- In our case $RO = \emptyset$, $WF = \emptyset$, $RW = \{j-1\}$
- Over loop DO j: $RO_{loop} = \emptyset$, $WF_{loop} = \emptyset$, $RW_{loop} = [0, NS-1]$



Summarizing Subroutine matmul

RW summary for matmul; $RO_{matmul} = WF_{matmul} = \emptyset$

```

 $S_{matmul}$   SUBROUTINE matmul(XE, NS)
            INTEGER NS, XE(*)
 $S_{loop}$     DO j = 1, NS
 $S_{RO}$       ... = XE(j) ...
 $S_{WF}$       XE(j) = ...
            ENDDO
            END
  
```

$$RW_{S_{loop}}^{XE} = [0, NS - 1]$$

$$S_{RO} \diamond S_{WF} = \{\emptyset, \emptyset, RW = \{j - 1\}\}$$

$$RO_{S_{RO}}^{XE} = \{j - 1\}$$

$$WF_{S_{WF}}^{XE} = \{j - 1\}$$

- Composing read-only RO_{S_1} and write-first WF_{S_2} regions:
- $RO = RO_{S_1} - WF_{S_2}$, $WF = WF_{S_2} - RO_{S_1}$, $RW = RO_{S_1} \cap WF_{S_2}$
- In our case $RO = \emptyset$, $WF = \emptyset$, $RW = \{j - 1\}$
- Over loop DO j: $RO_{loop} = \emptyset$, $WF_{loop} = \emptyset$, $RW_{loop} = [0, NS - 1]$



Summarizing Accesses for the Target Loop

RW summary for loop D0 i: $RW^i = ?$


```

Sloop   INTEGER NS, NP, IA(*), XE(*)
SWF    DO i = 1, N
SWF      CALL geteu (XE(IA(i)),NP,SYM)
SRW      CALL matmul(XE(IA(i)),NS)
          ENDDO

```

$S_{WF} \diamond S_{RW} = \{\emptyset, WF^i = WF_{geteu}, RW^i\}$

$RW^i =$



In our case, a sufficient condition for XE independence is:

$$\bigcup_{i=1}^N (RW_i \cap (\bigcup_{k=1}^{i-1} RW_k)) = \emptyset \quad \Leftrightarrow \quad RW_i = \emptyset \quad \Leftrightarrow$$

$$\Leftrightarrow \quad \text{SYM} \neq 1 \wedge \text{NS} \leq 16 * \text{NP}$$



Summary-Based Independence Equations

Flow and Anti Independence Equation for loop of index i :

$$\begin{aligned}
 S_{find} = & \{(\cup_{i=1}^N WF_i) \cap (\cup_{i=1}^N RO_i)\} \cup \\
 & \{(\cup_{i=1}^N WF_i) \cap (\cup_{i=1}^N RW_i)\} \cup \\
 & \{(\cup_{i=1}^N RO_i) \cap (\cup_{i=1}^N RW_i)\} \cup \\
 & \{\cup_{i=1}^N (RW_i \cap (\cup_{k=1}^{i-1} RW_k))\} = \emptyset
 \end{aligned} \tag{2}$$

Output Independence Equation for loop of index i :

$$S_{oind} = \{\cup_{i=1}^N (WF_i \cap (\cup_{k=1}^{i-1} WF_k))\} = \emptyset \tag{3}$$

Computing S_{find} and S_{oind} solves a more difficult problem than we need, i.e., computes the indexes involved in cross-iteration deps.

Loop Independence: when are S_{find} and S_{oind} empty?




Key Idea: Predicate-Centric Approach

Approach centered on extracting arbitrarily-shaped predicates.

Key Idea

- *Source of inaccuracy*: summary representation not closed under composition w.r.t. set operations.
- *Language* representation for summaries ... precise but *expensive* to compute at runtime
- "Let's *reason* about it!"
 $8 * NP < NS + 6 \Rightarrow A - B = \emptyset \Rightarrow S = \emptyset!$

Flow Independence USR Example:



// Show Calculix from SPEC2006!

A Not So Trivial Loop

Key Idea

$k = k0$	<i>Ind.</i>	$k = k0$
DO $i=0, N-1$	<i>Var.</i>	DO $i = 0, N-1$
$k = k+2$	\Rightarrow	$a(k0+2*i)=..$
$a(k)=..$	<i>Sub.</i>	ENDDO
ENDDO		$k=k0+MAX(2N-2,0)$
(a)		(b)

```

DO i = 0, N-1
  IF(cond(b(i)))THEN
    civ = civ + 1  $\Rightarrow$  ?
    a(civ) = ...
  ENDIF ENDDO
(c)

```

- After induction variable substitution, loop (b) is easily found parallel: $k0 + 2*i_1 = k0 + 2*i_2 \Rightarrow i_1 = i_2$



A Not So Trivial Loop

Key Idea

$k = k0$	<i>Ind.</i>	$k = k0$
DO $i=0, N-1$	<i>Var.</i>	DO $i = 0, N-1$
$k = k+2$	\Rightarrow	$a(k0+2*i)=..$
$a(k)=..$	<i>Sub.</i>	ENDDO
ENDDO		$k=k0+MAX(2N-2,0)$
(a)		(b)

```

DO i = 0, N-1
  IF(cond(b(i)))THEN
    civ = civ + 1  $\Rightarrow$  ?
    a(civ) = ...
  ENDIF ENDDO
(c)

```

- After induction variable substitution, loop (b) is easily found parallel: $k0 + 2*i_1 = k0 + 2*i_2 \Rightarrow i_1 = i_2$
- Loop (c) is similar to (a) but it is more difficult to reason about because *civ* cannot be expressed in terms of the loop index.
- Denote by civ_0 , civ_μ^i , civ_γ^i the values of *civ* at the beginning of the loop, and at the start and end of iteration *i*.
- Express the iteration summary on each path:
 - THEN branch: $\{civ_\mu^i+1\} = [civ_\mu^i+1, civ_\gamma^i]$
 - ELSE branch: $\emptyset = [civ_\mu^i+1, civ_\gamma^i]$, (an empty interval has its lower bounds $<$ its upper bound.)



A Not So Trivial Loop

Key Idea

$k = k0$	<i>Ind.</i>	$k = k0$
DO $i=0, N-1$	<i>Var.</i>	DO $i = 0, N-1$
$k = k+2$	\Rightarrow	$a(k0+2*i)=..$
$a(k)=..$	<i>Sub.</i>	ENDDO
ENDDO		$k=k0+MAX(2N-2,0)$
(a)		(b)

```

DO i = 0, N-1
  IF(cond(b(i)))THEN
    civ = civ + 1  $\Rightarrow$  ?
    a(civ) = ...
  ENDIF
ENDDO
(c)

```

- We have deduced that $W_i = [civ_{\mu}^i + 1, civ_{\gamma}^i]$
- Summarize the accesses of the first $i - 1$ iterations:
 - $civ_{\mu}^{i+1} = civ_{\gamma}^{i-1}$ since the value of civ at the start of an iteration is the same as the one at the end of previous iteration.
 - $\cup_{k=0}^{i-1} W_k = [civ_{\mu}^0 + 1, civ_{\gamma}^0] \cup \dots \cup [civ_{\mu}^{i-1} + 1, civ_{\gamma}^{i-1}] = [civ_{\mu}^0 + 1, civ_{\gamma}^{i-1}] = [civ_{\mu}^0 + 1, civ_{\mu}^i]$
- $(\cup_{k=0}^{i-1} W_k) \cap W_i = [civ_{\mu}^0 + 1, civ_{\mu}^i] \cap [civ_{\mu}^i + 1, civ_{\gamma}^i] = \emptyset$
- We have just proved NO output dependencies occur on array a!



A Not So Trivial Loop

Solving the True-Dependence (RAW) on `civ`

```
DO i = 0, N-1
  IF(cond(b(i)))THEN
    civ = civ + 1  $\Rightarrow$  ?
    a(civ) = ...
  ENDIF ENDDO
```

- Extract the loop slice that computes the `civ` value. If all accesses to `civ` are in reduction statement then compute partial contributions of each iteration.
- exclusive scan on the per-iteration contributions give the value of `civ` at the beginning of each iteration, and
- is plugged in the original loop.

```
civ0 = civ
DOALL i = 0, N-1
  civ = 0
  IF(cond(b(i))) THEN
    civ = civ + 1
  ENDIF
  civs[i] = civ
ENDDO

civs' = scanExc (+) 0 civs

DOALL i = 0, N-1
  civ = civ0 + civs'[i]
  IF(cond(b(i))) THEN
    civ = civ + 1
    a(civ) = ...
  ENDIF
ENDDO
```

